

# **Tutorial 10: Concurrency Control**

**CS3402 Database Systems**

# Question 1

- Consider the following table scheme, and assume that  $R_1$  has 1000 tuples,  $R_2$  has 1500 tuples and  $R_3$  has 750 tuples.
- $R_1 (\underline{A}, B, C)$
  - $R_2 (\underline{C}, D, E)$
  - $R_3 (\underline{E}, F)$
- a) Estimate the size (max and min numbers of tuples) of  $R_1 * R_2 * R_3$  (where  $*$  denotes Natural Join).
- b) There are two ways to perform the Natural Join in (a), which one is more efficient in terms of number of comparison?
- $(R_1 * R_2) * R_3$
  - $R_1 * (R_2 * R_3)$

# Question 1a (Answer)

- $R_1$  has 1,000 tuples, so A has 1,000 distinct values
- $R_2$  has 1,500 tuples, so C has 1,500 distinct values
- $R_3$  has 750 tuples, so E has 750 distinct values
- $\text{Temp} \leftarrow R_1 * R_2$  produces 0 to 1,000 tuples.
- $\text{Temp} * R_3$  produces 0 to 1,000 tuples.
- The min and max numbers of tuples are 0 and 1,000, respectively.

## Question 1b (Answer) (1/2)

- $R_1$  has 1,000 tuples, so A has 1,000 distinct values
- $R_2$  has 1,500 tuples, so C has 1,500 distinct values
- $R_3$  has 750 tuples, so E has 750 distinct values
- For (i),  $(R_1 * R_2) * R_3$ 
  - $\text{Temp} \leftarrow R_1 * R_2$  requires  $1,000 * 1,500 = 1,500,000$  comparisons at the worst case
  - $\text{Temp} * R_3$  requires  $1,000 * 750 = 750,000$  comparisons at the worst case
  - In total, it requires 2,250,000 comparisons at the worst case

## Question 1b (Answer) (2/2)

- $R_1$  has 1,000 tuples, so A has 1,000 distinct values
- $R_2$  has 1,500 tuples, so C has 1,500 distinct values
- $R_3$  has 750 tuples, so E has 750 distinct values
- For (ii),  $R_1 * (R_2 * R_3)$ 
  - $\text{Temp} \leftarrow R_2 * R_3$  requires  $1,500 * 750 = 1,125,000$  comparisons at the worst case
  - $R_1 * \text{Temp}$  requires  $1,000 * 1,500 = 1,500,000$  comparisons at the worst case
  - In total, it requires 2,625,000 comparisons at the worst case
- In conclusion, (i) is more efficient than (ii) in terms of number of comparison.

## Question 2 (1/2)

- A canonical query tree is a tree structure that corresponds to a relational algebra expression or an SQL query directly, without doing any optimization. As such, it is usually not the most efficient way of executing the query.
- Consider the relations:

EMPLOYEE(ENAME, SSN, BDATE, ADDRESS, DNUM)

PROJECT(PNAME, PNUMBER, PLOCATION, DNUM)

WORKS\_ON(ESSN, PNO, HOURS)

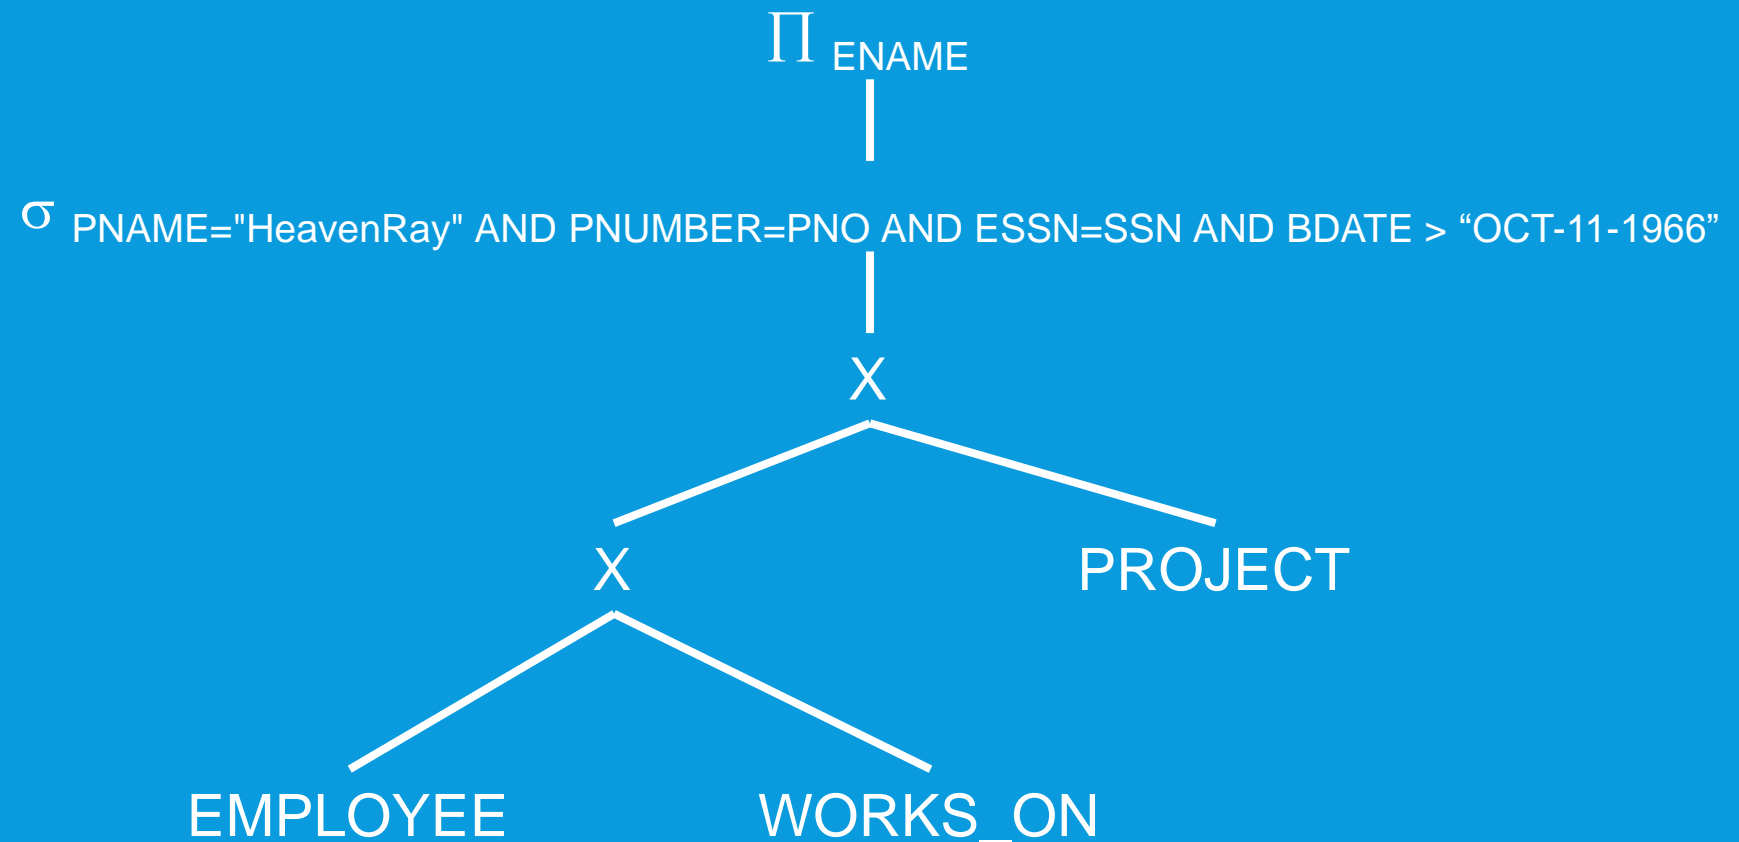
## Question 2 (2/2)

- And the following SQL query:

```
SELECT      ENAME
FROM        EMPLOYEE, WORKS_ON, PROJECT
WHERE       PNAME="HeavenRay" AND PNUMBER=PNO
           AND ESSN=SSN AND BDATE > "OCT-11-1966";
```

- a) Draw a canonical query tree for the above SQL query.
- b) Apply the optimization rules to the above query tree and come up with the most optimized query tree.

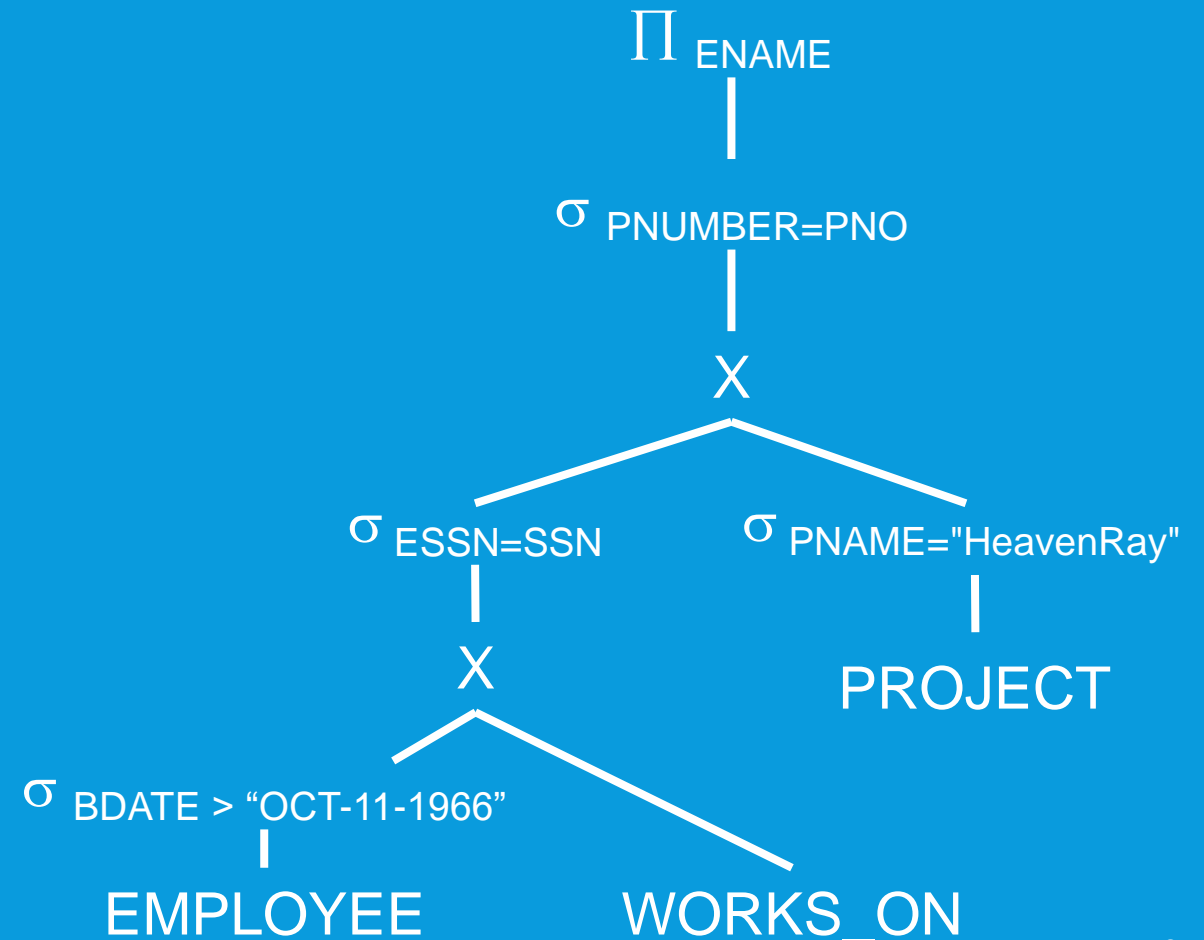
## Question 2a (Answer) (1/5)





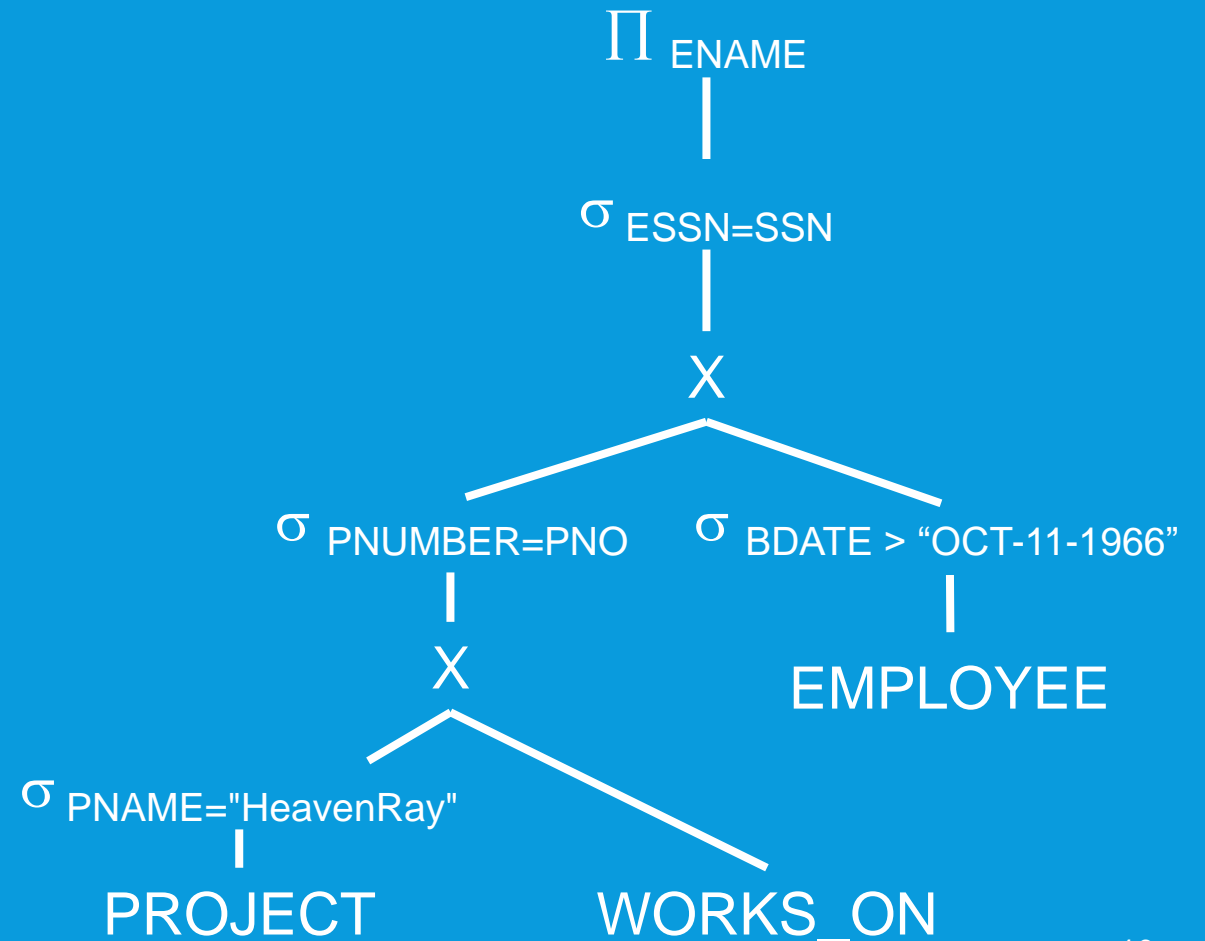
## Question 2b (Answer) (2/5)

- Steps 1 and 2: Break up any SELECT operations with conjunctive conditions and move each SELECT operation as far down the query tree



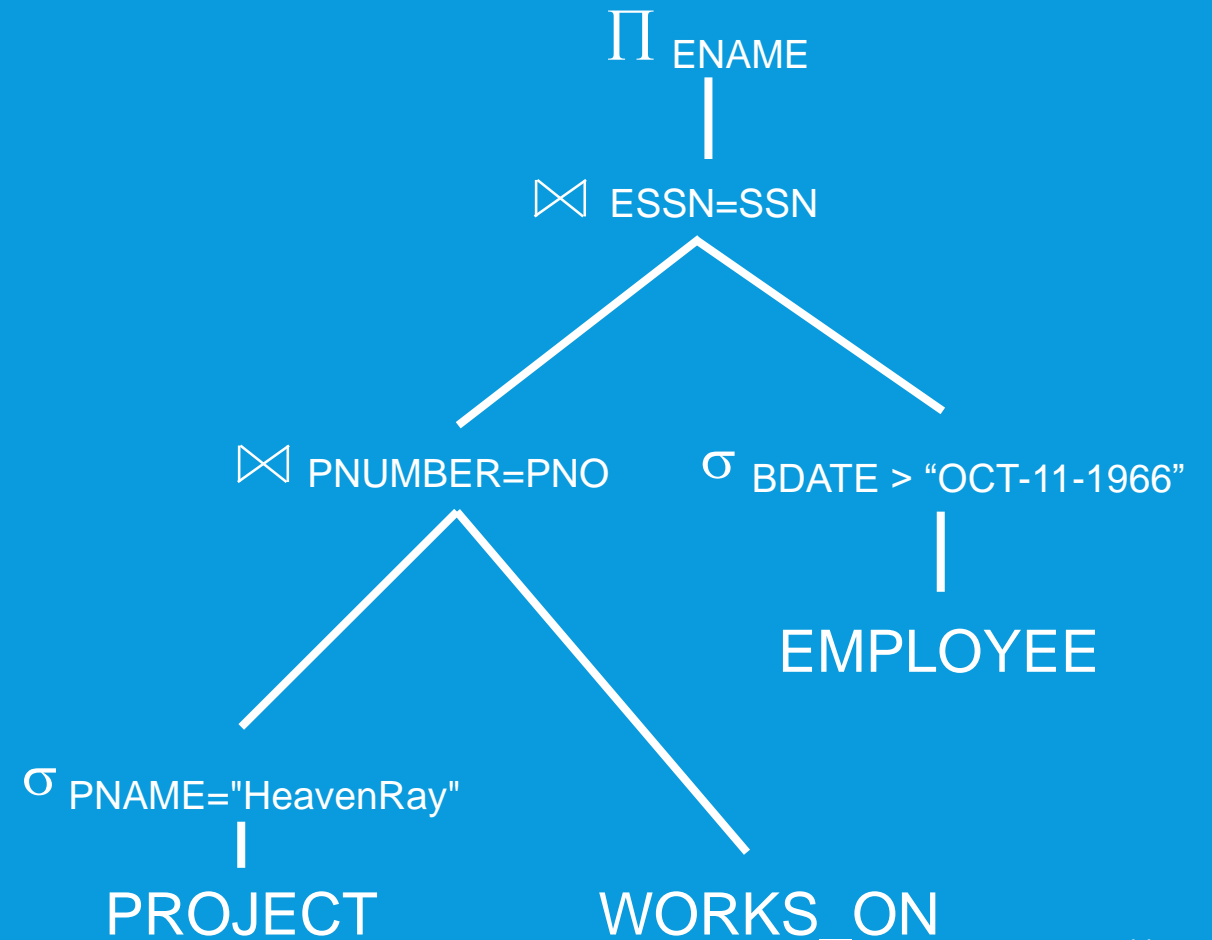
## Question 2b (Answer) (3/5)

- Step 3: Applying the more restrictive SELECT operation first



## Question 2b (Answer) (4/5)

- Step 4: Replacing CROSS PRODUCT and SELECT with JOIN operation



## Question 2b (Answer) (5/5)

- Step 5: Moving PROJECT operations down the query tree

