

**Question 1.** Suppose that we construct a B<sup>+</sup> tree with 3 levels, in which the order of internal nodes is 20 and the order of leaf nodes is 20, what is the maximum number of data record pointers for this B<sup>+</sup> tree? Please show the number of key entries (or data record entries) and the number of pointers at the root, level 1, level 2 and leaf level. [20 marks]

Root:	1 node	19 key entries	20 pointers	[5 marks]
Level 1:	20 nodes	380 key entries (20*19)	400 pointers	[5 marks]
Level 2:	400 nodes	7,600 key entries (400*19)	8,000 pointers	[5 marks]
Leaf level:	8,000 pointers	152,000 data record entries (8,000*19)  Alternative solution: 160,000 data record entries (8,000*20)		[5 marks]

The maximum number of data record pointers is 152,000 or 160,000.

**Question 2.** Suppose that the size of a search key field is V=9 bytes, the size of a record pointer is Pr=7 bytes, the size of a block pointer/tree pointer is P=6 bytes, the order of internal nodes (p) is 20, and the number of data record points for a leaf node 20 for a B<sup>+</sup> tree. What should be the minimum required block size B for this B<sup>+</sup> tree. [20 marks]

For internal nodes:

$$(p * P) + ((p - 1) * V) \leq B$$

$$(20 * 6) + ((20 - 1) * 9) \leq B$$

$$291 \leq B \text{ [8 marks]}$$

For leaf nodes:

$$(p_{\text{leaf}} * (Pr + V)) + P \leq B$$

$$(20 * (7 + 9)) + 6 \leq B$$

$$326 \leq B \text{ [8 marks]}$$

The minimum required block size B is 326. [4 marks]

Course: CS3402 Database Systems  
Term: 2020-21 Semester B  
Assignment #3 (100 marks)  
Submission deadline: May 01, 2021

Name: \_\_\_\_\_  
Student ID: \_\_\_\_\_

**Question 3.** The following table shows the schedule for transactions  $T_1$ ,  $T_2$  and  $T_3$  with  $T_1$  having an “older” timestamp than  $T_2$  and  $T_2$  is older than  $T_3$ . Use `unlock(A)` and `unlock(B)` to release a `read_lock` or `write_lock` on A and B, respectively.

$T_1$	$T_2$	$T_3$
read(A);		
	read(B);	
write(B);		
		write(A);
	write(A);	
commit;		
	commit;	
		write(B);
		commit;

(a) If B2PL (Basic Two-Phase Locking) is used for concurrency control, does B2PL cause any deadlock? If yes, what is the cycle in the wait-for graph. If no, what is the final schedule for these three transactions? [20 marks]

[This table is just for reference; it is not part of the answer.]

T <sub>1</sub>	T <sub>2</sub>	T <sub>3</sub>
read_lock(A); read(A);		
	read_lock(B); read(B);	
write(B); (blocked)		
		write(A); (blocked)
	write(A); (blocked)	
commit;		
	commit;	
		write(B);
		commit;

Yes, B2PL causes a deadlock. [5 marks]

The cycle in the wait-for-graph is  $T_1 \rightarrow T_2 \rightarrow T_1$  (or  $T_2 \rightarrow T_1 \rightarrow T_2$ ). [15 marks]

(b) If C2PL (Conservative Two-Phase Locking) is used for concurrency control, does C2PL cause any deadlock? If yes, what is the cycle in the wait-for graph. If no, what is the final schedule for these three transactions? [20 marks]

[This table is just for reference; it is not part of the answer.]

T <sub>1</sub>	T <sub>2</sub>	T <sub>3</sub>
read_lock(A) ; write_lock(B);		
read(A);		
write(B);		
commit;		
unlock(A); unlock(B);		
	read_lock(B); write_lock(A);	
	read(B);	
	write(A);	
	commit;	
	unlock(B);unlock(A);	
		write_lock(A); write_lock(B);
		write(A);
		write(B);
		commit;
		unlock(A); unlock(B);

No, C2PL does not cause any deadlock. [5 marks]

The final schedule is T<sub>1</sub>, T<sub>2</sub> and T<sub>3</sub> (or T<sub>1</sub> => T<sub>2</sub> => T<sub>3</sub>). [15 marks]

(c) If the S2PL (Strict Two-Phase Locking) is used for concurrency control and the wait-die method is used to prevent deadlock, which transaction(s) will be restarted? For a transaction that must be restarted, please state the reasons (including which transaction and database item it restarts for). [20 marks]

[This table is just for reference; it is not part of the answer.]

T <sub>1</sub>	T <sub>2</sub>	T <sub>3</sub>
read_lock(A); read(A);		
	read_lock(B); read(B);	
write(B); (blocked)		
		write(A); (T <sub>3</sub> restarts because T <sub>3</sub> is younger than T <sub>1</sub> .)
	write(A); (T <sub>2</sub> restarts because T <sub>2</sub> is younger than T <sub>1</sub> .)	
write_lock(B); write(B);		
commit;		
unlock(A); unlock(B);		

T<sub>3</sub> will be restarted. T<sub>3</sub> waits for database item A which is being locked by T<sub>1</sub>. [10 marks]

T<sub>2</sub> will be restarted. T<sub>2</sub> waits for database item A which is being locked by T<sub>1</sub>. [10 marks]