


# Transport Layer



*Computer  
Networking: A Top  
Down Approach*  
7<sup>th</sup> edition  
Jim Kurose, Keith Ross  
Pearson, 2017

# Intended Learning Outcomes

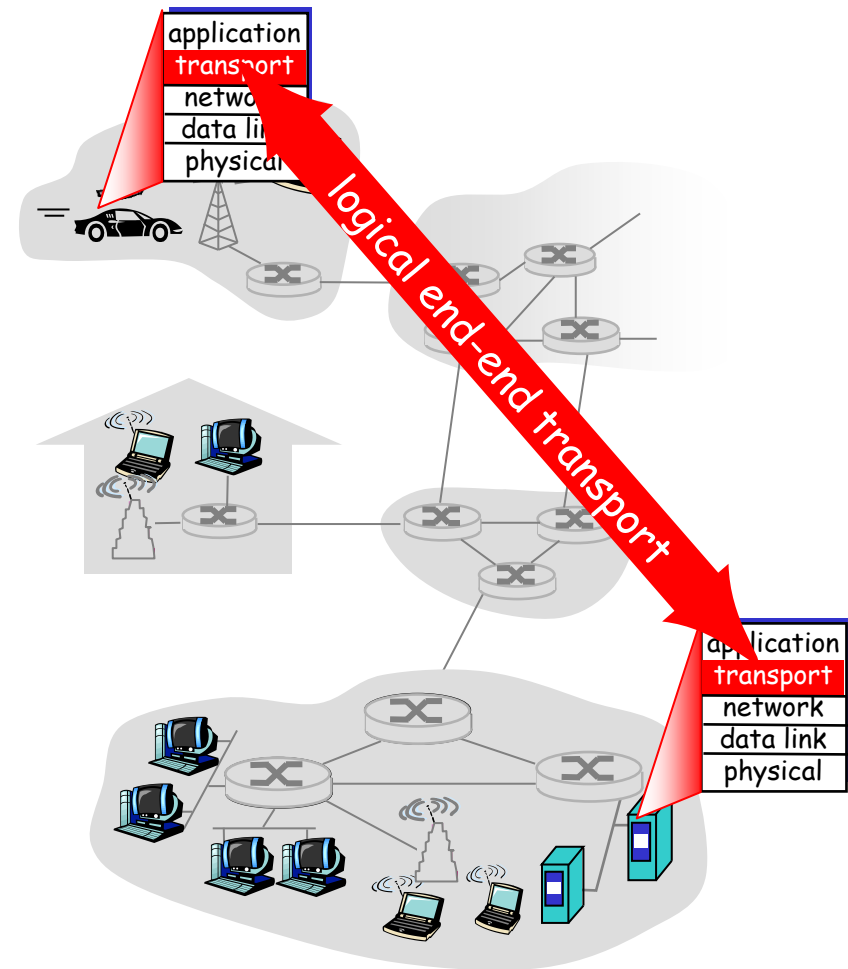
- ❑ understand principles behind transport layer services:
  - reliable data transfer
  - flow control
  - congestion control
- ❑ understand transport layer protocols in the Internet:
  - TCP: connection-oriented transport, connection management
  - TCP flow control
  - TCP congestion control

# Transport Layer

- ❑ Transport-layer services
- ❑ Connectionless transport: UDP
- ❑ Connection-oriented transport: TCP
  - segment structure
  - reliable data transfer
  - flow control
  - connection management
- ❑ Principles of congestion control
- ❑ TCP congestion control

# Transport services and protocols

- ❑ provide *logical communication* between app processes running on different hosts
- ❑ transport protocols run in end systems
  - send side: breaks app messages into *segments*, passes to network layer
  - rcv side: reassembles segments into messages, passes to app layer
- ❑ more than one transport protocol available to apps
  - Internet: TCP and UDP



# Transport Layer

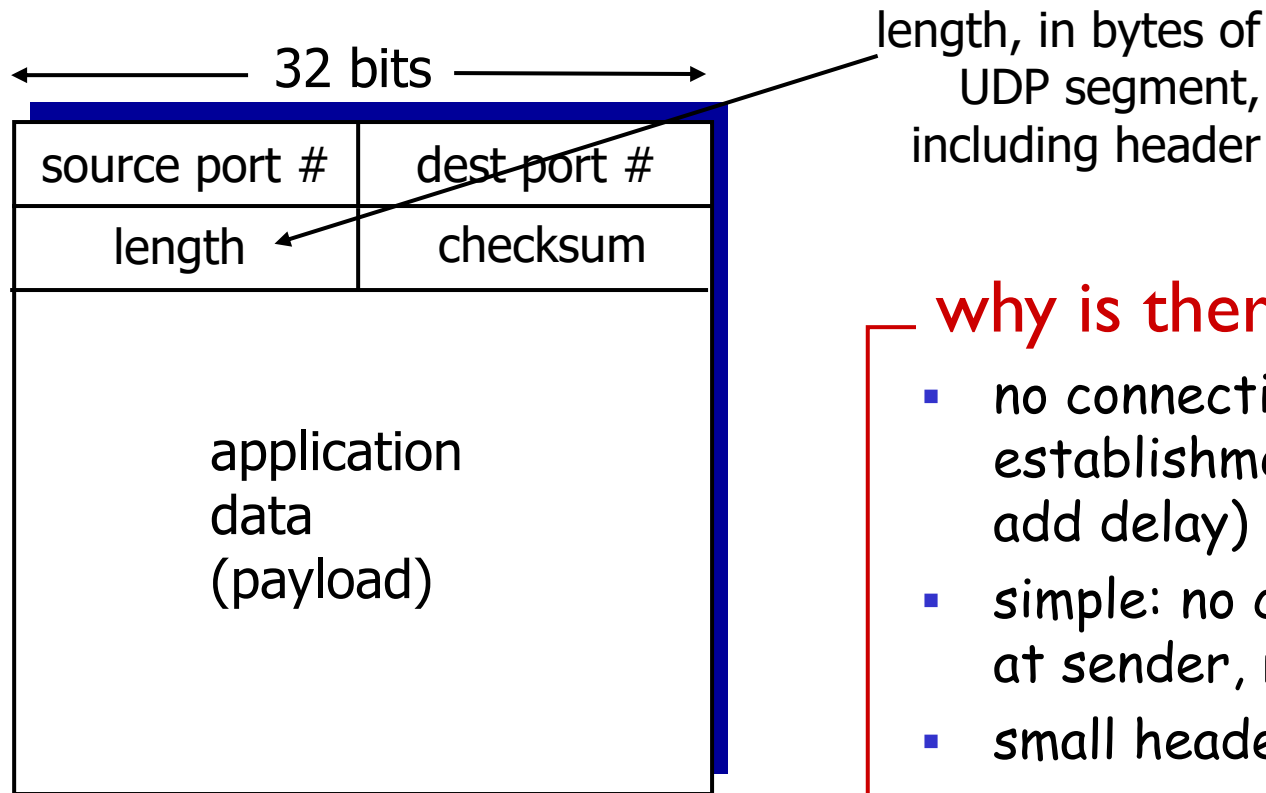
- ❑ Transport-layer services
- ❑ Connectionless transport: UDP
- ❑ Connection-oriented transport: TCP
  - segment structure
  - reliable data transfer
  - flow control
  - connection management
- ❑ Principles of congestion control
- ❑ TCP congestion control

# UDP: User Datagram Protocol [RFC 768]

- ❑ “no frills,” “bare bones” Internet transport protocol
- ❑ “best effort” service, UDP segments may be:
  - lost
  - delivered out-of-order to app
- ❑ *connectionless*:
  - no handshaking between UDP sender, receiver
  - each UDP segment handled independently of others

- ❑ UDP use:
  - streaming multimedia apps (loss tolerant, rate sensitive)
  - DNS
  - SNMP
- ❑ reliable transfer over UDP:
  - add reliability at application layer
  - application-specific error recovery!

# UDP: segment header



UDP segment format

## — why is there a UDP? —

- no connection establishment (which can add delay)
- simple: no connection state at sender, receiver
- small header size
- no congestion control: UDP can send out data as fast as desired

# Transport Layer


- ❑ Transport-layer services
- ❑ Connectionless transport: UDP
- ❑ **Connection-oriented transport: TCP**
  - segment structure
  - reliable data transfer
  - flow control
  - connection management
- ❑ Principles of congestion control
- ❑ TCP congestion control



# TCP: Review

RFCs: 793, 1122, 1323, 2018, 2581

- ❑ **point-to-point:**
  - one sender, one receiver
- ❑ **reliable, in-order byte stream:**
  - no "message boundaries"
- ❑ **pipelined:**
  - TCP congestion and flow control set window size
- ❑ **send & receive buffers**

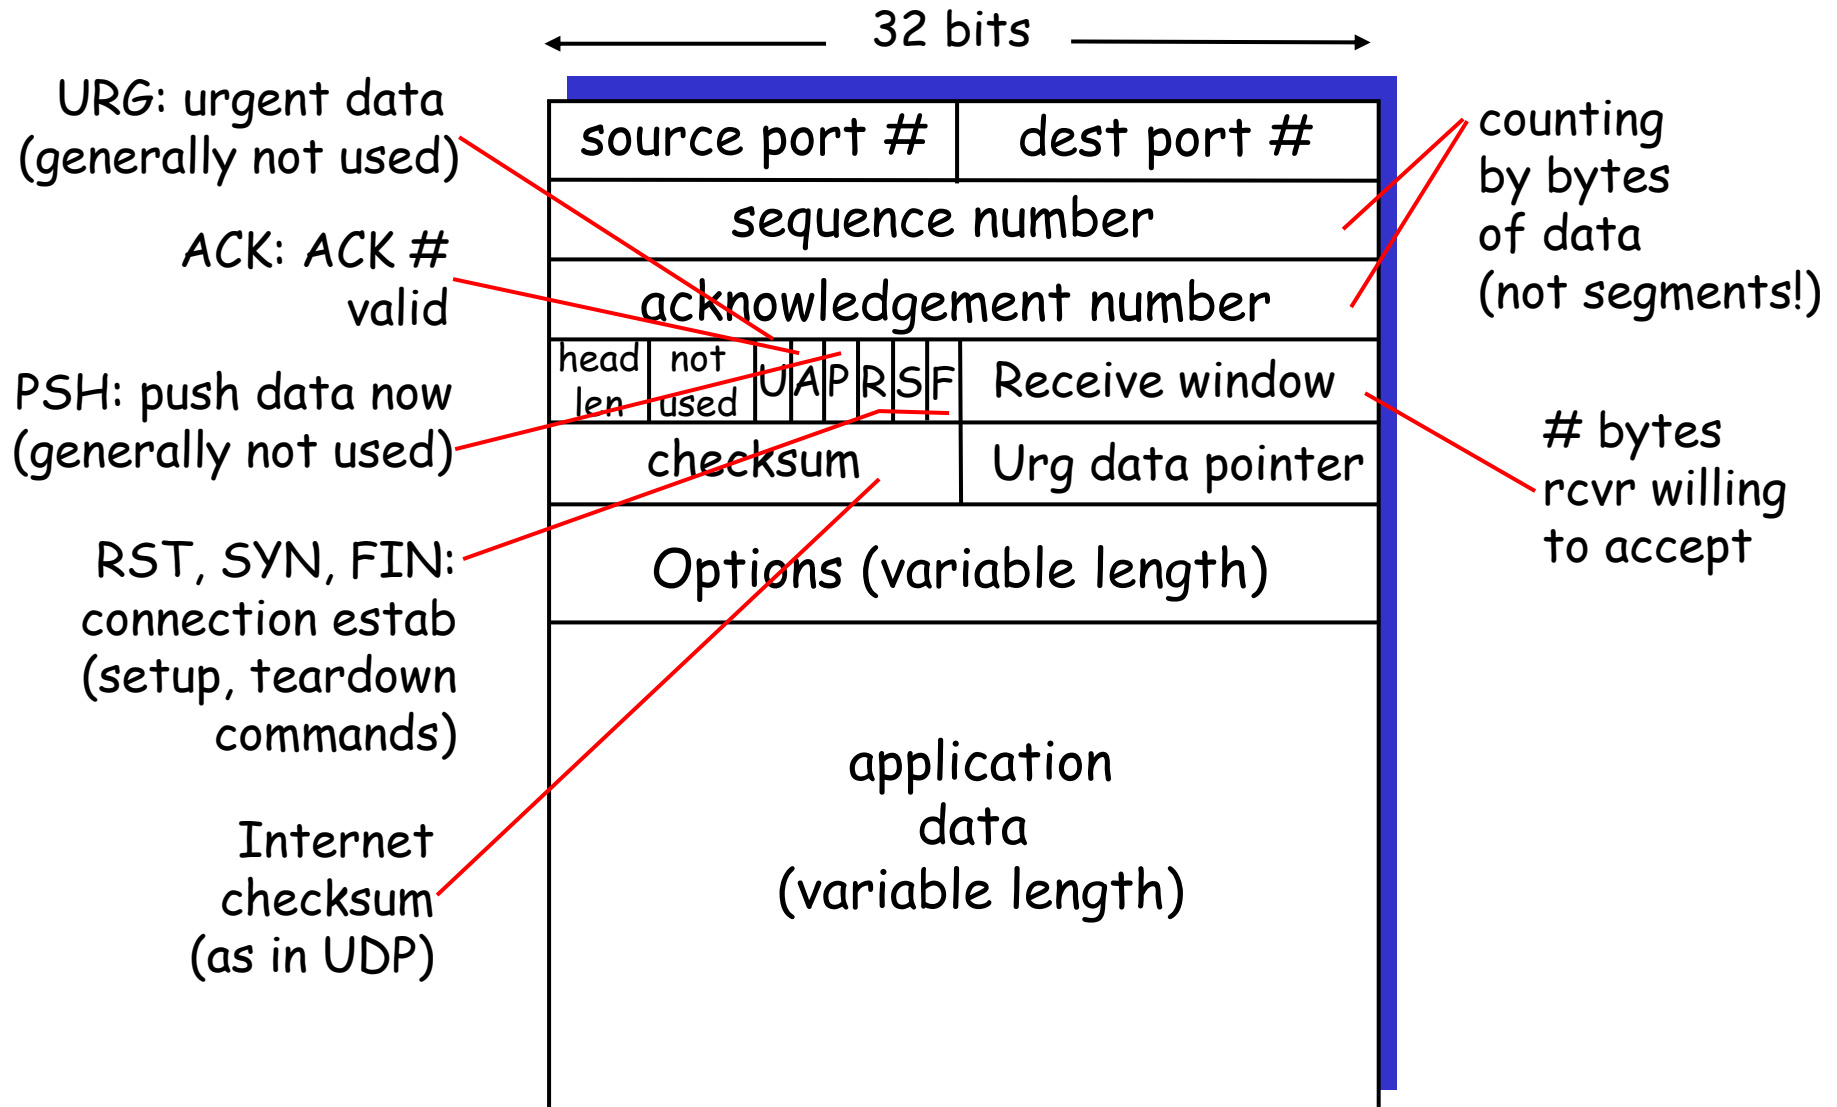


- ❑ **full duplex data:**
  - bi-directional data flow in same connection
  - MSS: maximum segment size
- ❑ **connection-oriented:**
  - handshaking (exchange of control msgs) init's sender & receiver states before data exchange
- ❑ **flow controlled:**
  - sender will not overload receiver
- ❑ **congestion controlled:**
  - sender will not overload network

# Transport Layer

- ❑ Transport-layer services
- ❑ Connectionless transport: UDP
- ❑ Connection-oriented transport: TCP
  - segment structure
  - reliable data transfer
  - flow control
  - connection management
- ❑ Principles of congestion control
- ❑ TCP congestion control

# TCP segment structure



# Transport Layer

- ❑ Transport-layer services
- ❑ Connectionless transport: UDP
- ❑ Connection-oriented transport: TCP
  - segment structure
  - reliable data transfer
  - flow control
  - connection management
- ❑ Principles of congestion control
- ❑ TCP congestion control

# TCP seq. numbers, ACKs

## sequence number:

- byte stream “number” of first byte in segment’s data

## acknowledgement number:

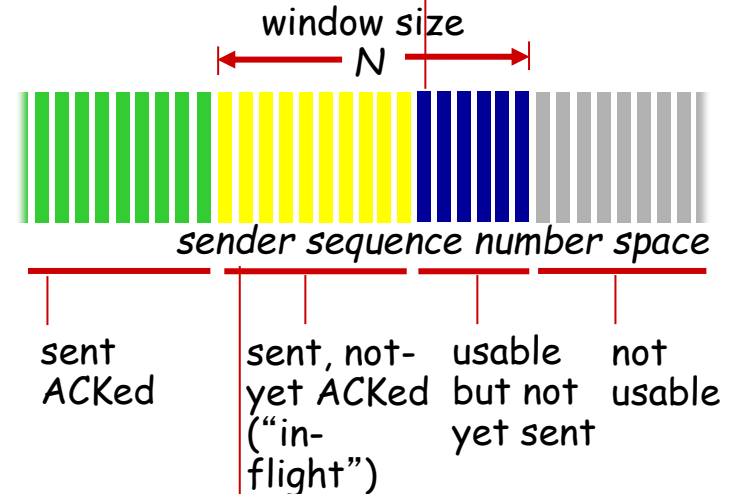
- seq # of next byte expected from other side
- cumulative ACK

**Q:** how receiver handles out-of-order segments

- Ans: TCP spec doesn’t say, - up to implementor

outgoing segment from sender

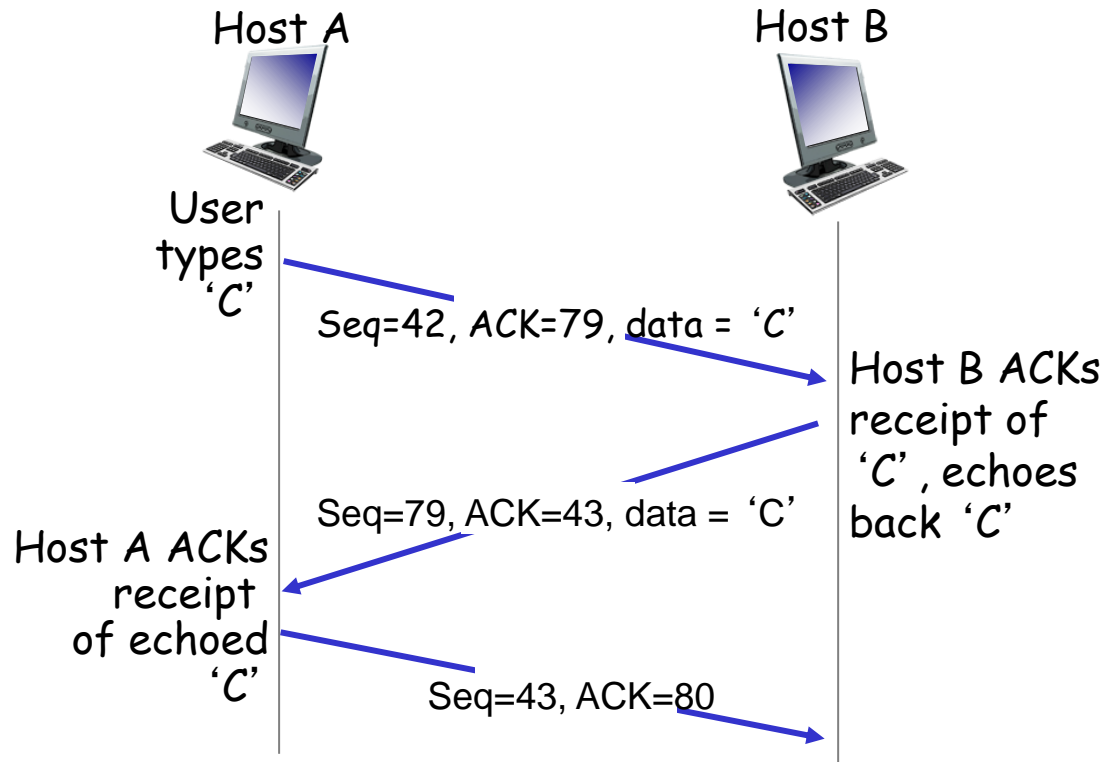
source port #	dest port #
sequence number	
acknowledgement number	
	rwnd
checksum	urg pointer



incoming segment to sender

source port #	dest port #
sequence number	
acknowledgement number	
	A
checksum	urg pointer

# TCP seq. numbers, ACK numbers



simple telnet scenario

# TCP reliable data transfer (rdt)

❖ TCP creates rdt service on top of IP's unreliable service

- pipelined segments
- cumulative acks
- single retransmission timer

❖ retransmissions triggered by:

- timeout events
- duplicate acks

let's initially consider simplified TCP sender:

- ignore duplicate acks
- ignore flow control, congestion control

# TCP sender events:

## *data rcvd from app:*

- ❖ create segment with seq #
- ❖ seq # is byte-stream number of first data byte in segment
- ❖ start timer if not already running
  - think of timer as for oldest unacked segment
  - expiration interval: `TimeoutInterval`

## *timeout:*

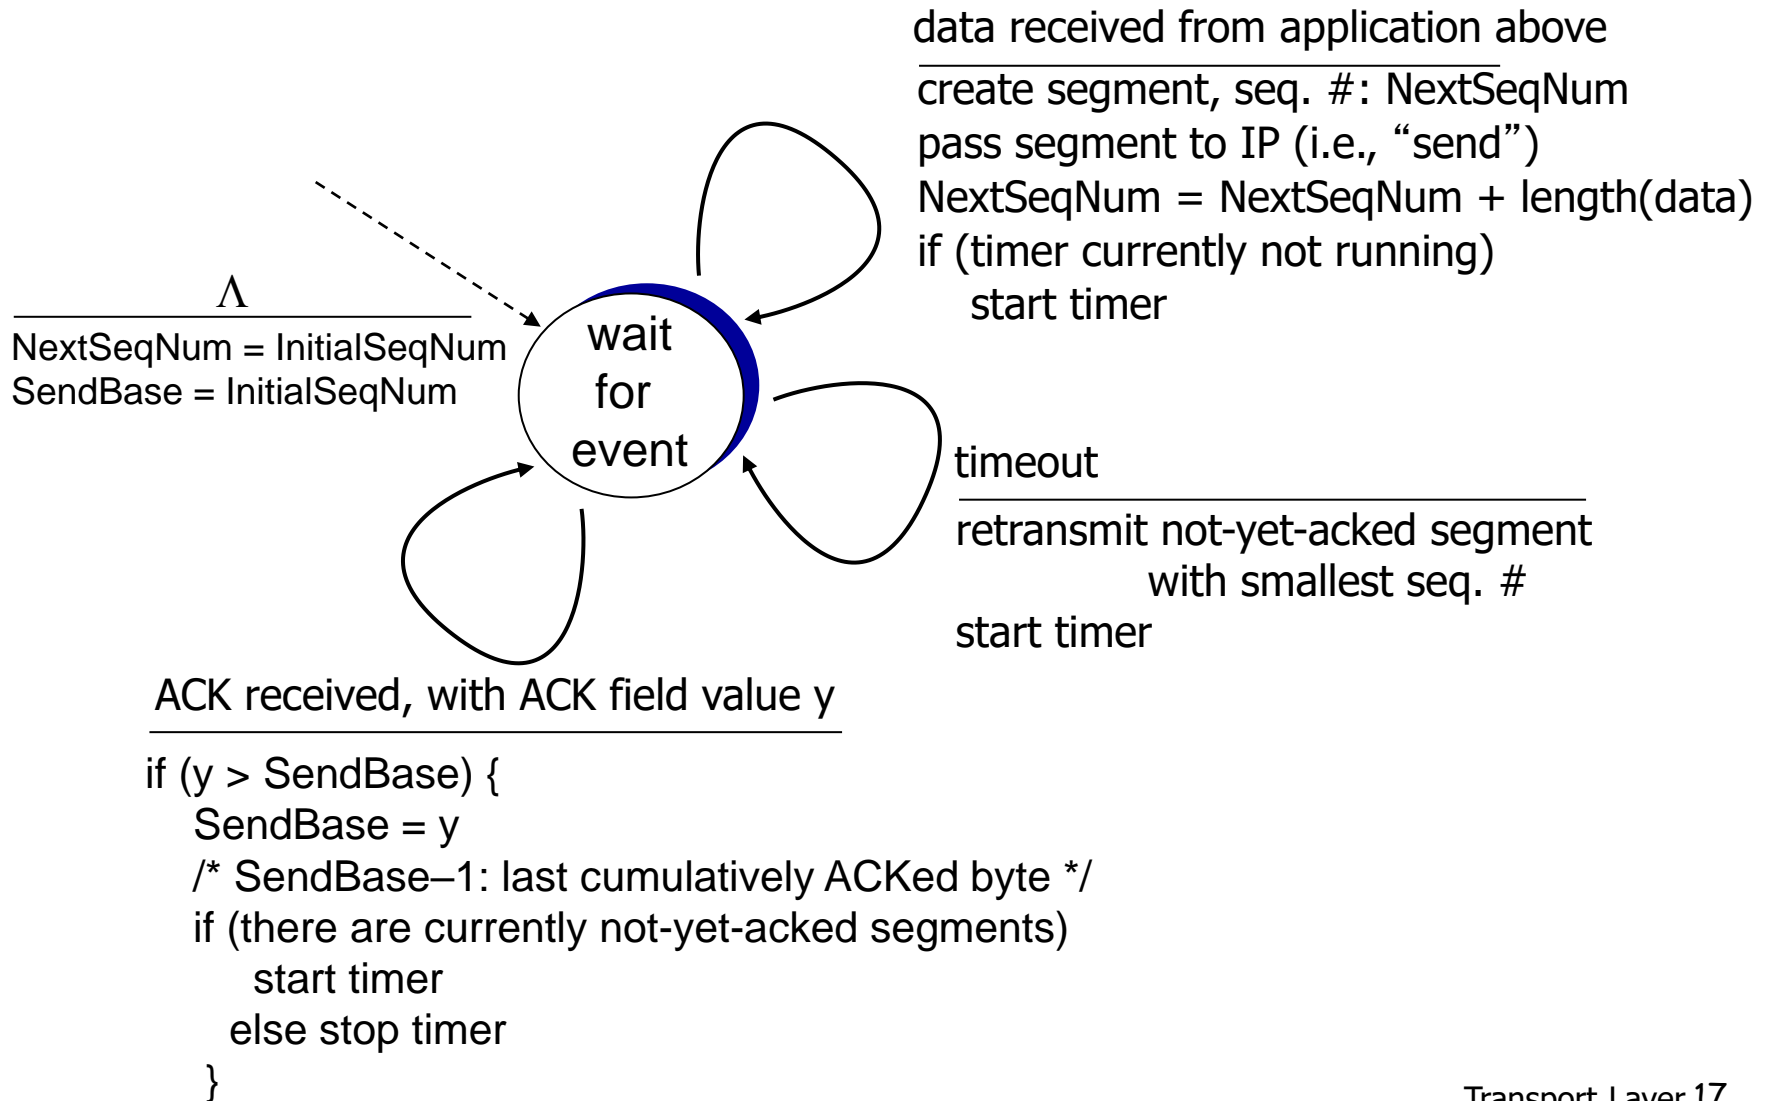
- ❖ retransmit segment that caused timeout
- ❖ restart timer

## *ack rcvd:*


- ❖ if ack acknowledges previously unacked segments
  - update what is known to be ACKed
  - start timer if there are still unacked segments




# TCP sender (simplified)



# TCP: retransmission scenarios




lost ACK scenario



premature timeout

# TCP: retransmission scenarios: Poll 6



# TCP ACK generation

## [RFC 1122, RFC 2581]

<i>event at receiver</i>	<i>TCP receiver action</i>
arrival of in-order segment with expected seq #. All data up to expected seq # already ACKed	delayed ACK. Wait up to 500ms for next segment. If no next segment, send ACK
arrival of in-order segment with expected seq #. One other segment has ACK pending	immediately send single cumulative ACK, ACKing both in-order segments
arrival of out-of-order segment higher-than-expect seq. # . Gap detected	immediately send <i>duplicate ACK</i> , indicating seq. # of next expected byte
arrival of segment that partially or completely fills gap	immediate send ACK, provided that segment starts at lower end of gap

# TCP fast retransmit

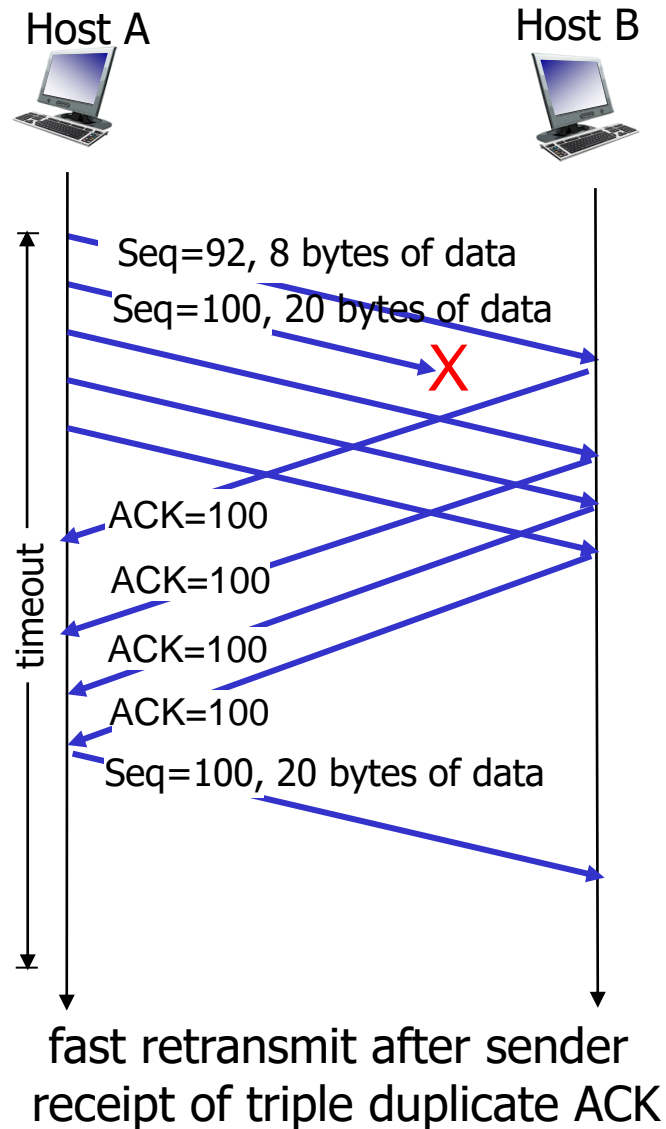
- ❖ time-out period often relatively long:
  - long delay before resending lost packet
- ❖ detect lost segments via duplicate ACKs.
  - sender often sends many segments back-to-back
  - if segment is lost, there will likely be many duplicate ACKs.

## *TCP fast retransmit*

if sender receives 4 ACKs for same data (“triple duplicate ACKs”), resend unacked segment with smallest seq #

- likely that unacked segment lost, so don't wait for timeout

# TCP fast retransmit



# Round trip time (RTT) vs timeout

Q: how to set TCP timeout value?

- ❖ longer than RTT
  - but RTT varies
- ❖ *too short*: premature timeout, unnecessary retransmissions
- ❖ *too long*: slow reaction to segment loss


Q: how to estimate RTT?

- ❖ **SampleRTT**: measured time from segment transmission until ACK receipt
  - ignore retransmissions
- ❖ **SampleRTT** will vary, want estimated RTT “smoother”
  - average several *recent* measurements, not just current **SampleRTT**

# Round trip time (RTT) vs timeout

$$\text{EstimatedRTT}_{i+1} = (1-\alpha) * \text{EstimatedRTT}_i + \alpha * \text{SampleRTT}_i$$

- ❖ exponential weighted moving average
- ❖ influence of past sample decreases exponentially fast
- ❖ typical value:  $\alpha = 0.125$





# Round trip time (RTT) vs timeout

- ❖ **timeout interval:** **EstimatedRTT** plus “safety margin”
  - larger variation in **EstimatedRTT** → larger safety margin
- ❖ estimate **SampleRTT** deviation from **EstimatedRTT**:

$$\text{DevRTT}_{i+1} = (1-\beta) * \text{DevRTT}_i + \beta * |\text{SampleRTT}_i - \text{EstimatedRTT}_i|$$

(typically,  $\beta = 0.25$ )

$$\text{TimeoutInterval}_i = \text{EstimatedRTT}_i + 4 * \text{DevRTT}_i$$



↑  
estimated RTT

↑  
“safety margin”

# Summary for TCP reliable data transfer

- ❑ How to make reliable data?
  - Sequence number, retransmission timer, cumulative ACK
- ❑ How to shorten retransmission delay?
  - "Fast Retransmit": lost segments detection via duplicate ACKs
- ❑ How to set time-out value?
  - Exponential weighted moving average

# Transport Layer

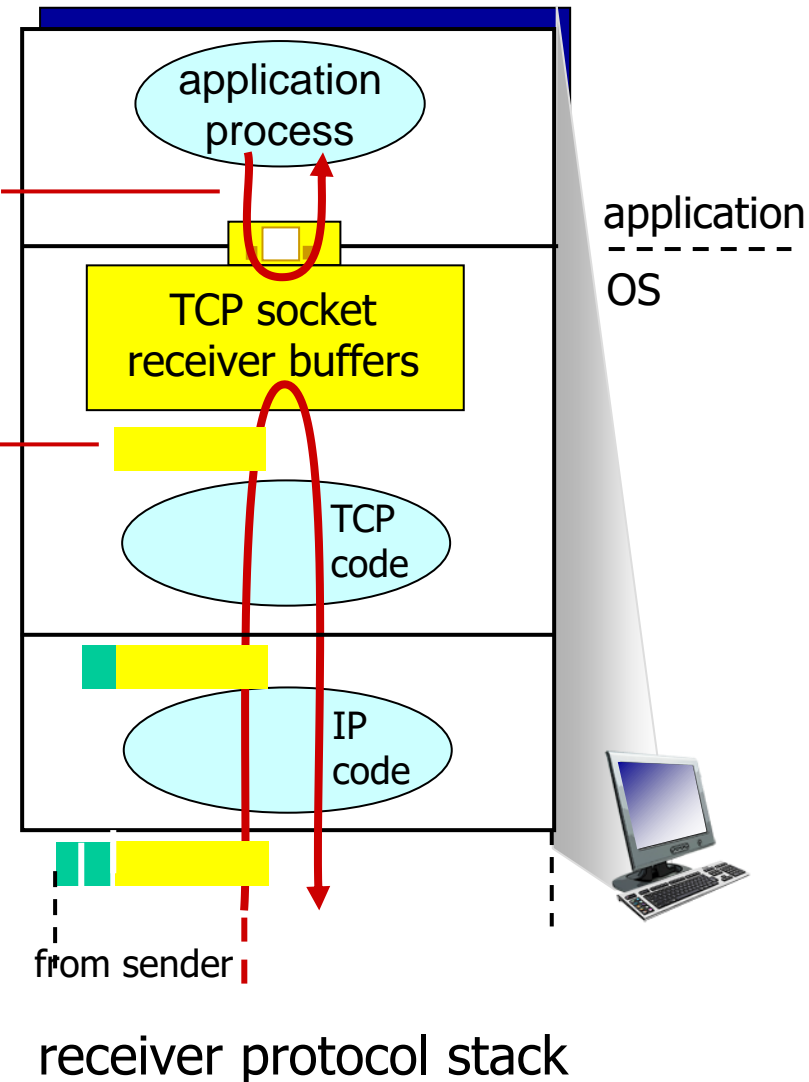
- ❑ Transport-layer services
- ❑ Connectionless transport: UDP
- ❑ Connection-oriented transport: TCP
  - segment structure
  - reliable data transfer
  - flow control
  - connection management
- ❑ Principles of congestion control
- ❑ TCP congestion control

# TCP flow control

application may  
remove data from  
TCP socket buffers ....

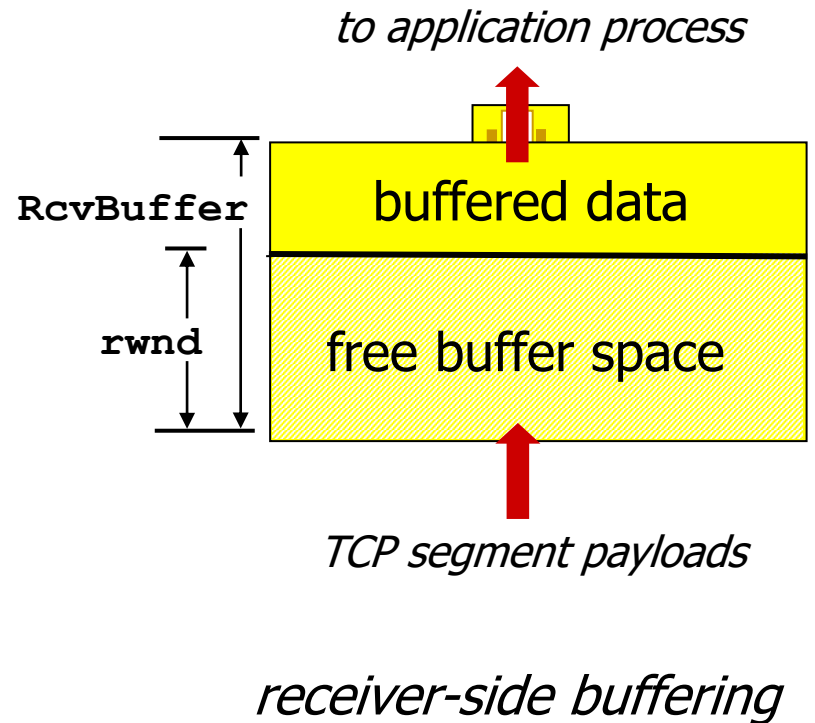
... slower than TCP  
receiver is delivering  
(sender is sending)

***flow control***  
receiver controls sender, so  
sender won't overflow  
receiver's buffer by transmitting  
too much, too fast



# TCP flow control (cont'ed)

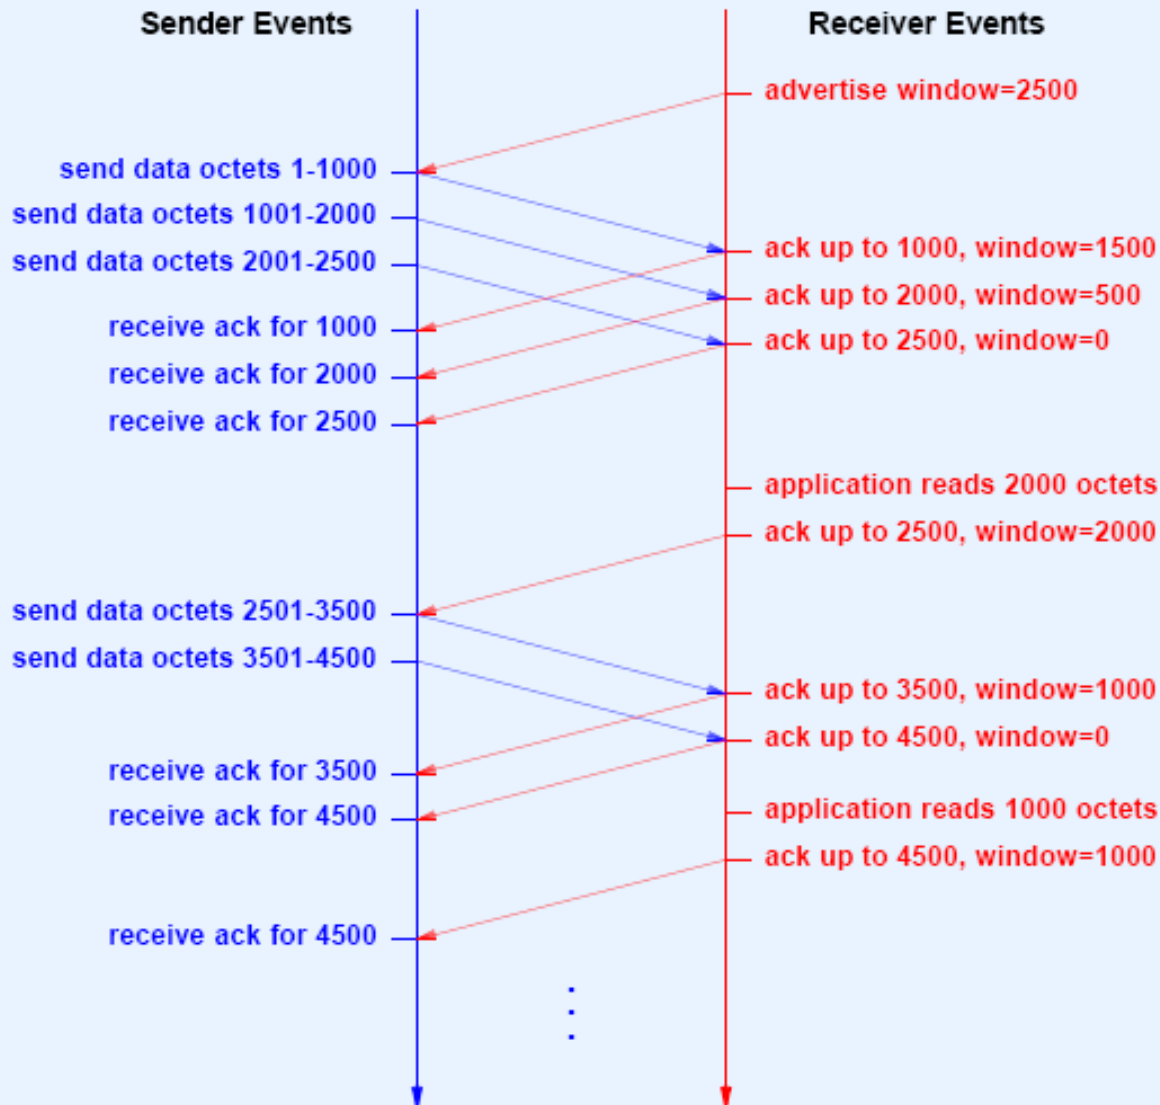
- ❖ receiver “advertises” free buffer space by including **rwnd** value in TCP header of receiver-to-sender segments
  - **RcvBuffer** size set via socket options (typical default is 4096 bytes)
  - many operating systems autoadjust **RcvBuffer**
- ❖ sender limits amount of unacked (“in-flight”) data to receiver’s **rwnd** value
- ❖ guarantees receiver buffer will not overflow



# Window Advertisement

- ❑ Each ACK carries new window information:
  - Acknowledgement number ( $AN$ )
  - Window size ( $W$ )
  
- ❑ ACK contains  $AN = i$ ,  $W = j$ :
  - Bytes through  $SN = i - 1$  acknowledged
    - Cumulative ACK
    - Byte  $i$  has not been received (It is the next byte expected)
  - Permission is granted to send  $W = j$  more bytes
    - i.e. bytes  $i$  through  $i + j - 1$

# Illustration: Window Advertisement



# Transport Layer


- ❑ Transport-layer services
- ❑ Connectionless transport: UDP
- ❑ Connection-oriented transport: TCP
  - segment structure
  - reliable data transfer
  - flow control
  - connection management
- ❑ Principles of congestion control
- ❑ TCP congestion control



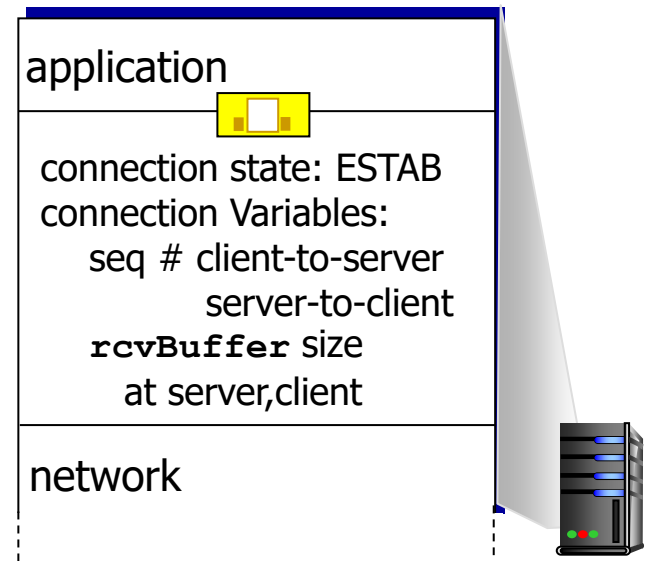
# Connection Management

before exchanging data, sender/receiver “handshake”:

- ❖ agree to establish connection (each knowing the other willing to establish connection)
- ❖ agree on connection parameters





```
Socket clientSocket =  
    newSocket("hostname", "port  
    number");
```



```
Socket connectionSocket =  
    welcomeSocket.accept();
```

# Agreeing to establish a connection

2-way handshake:



Q: will 2-way handshake always work in network?

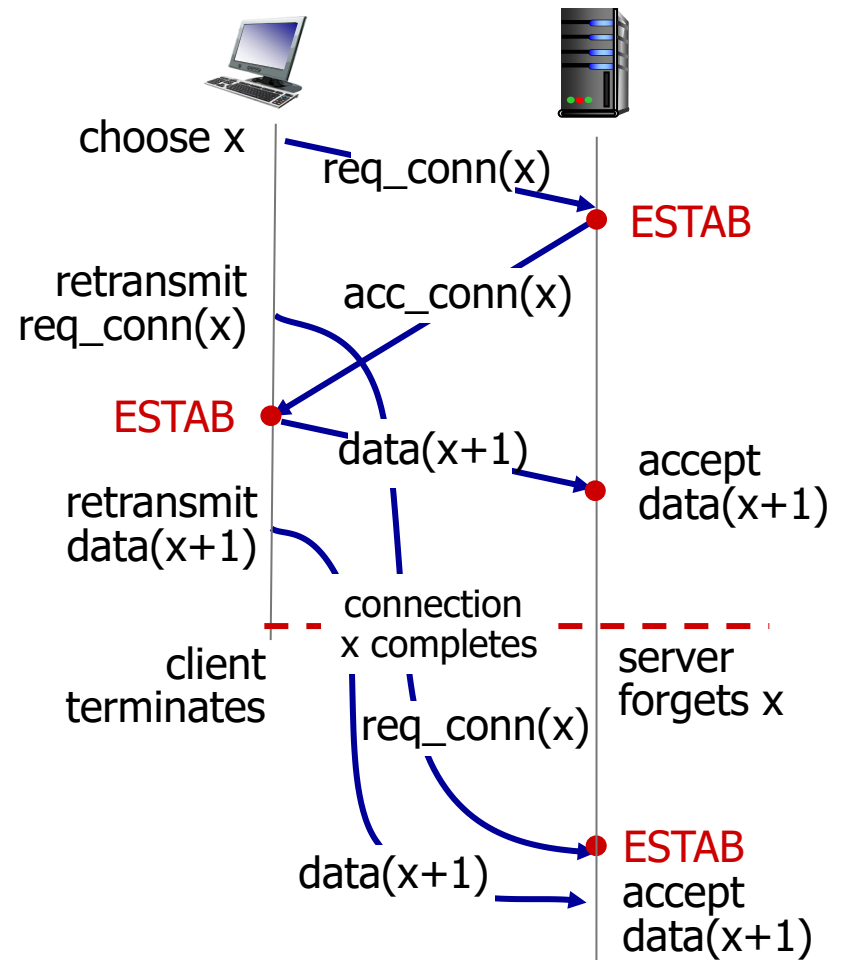
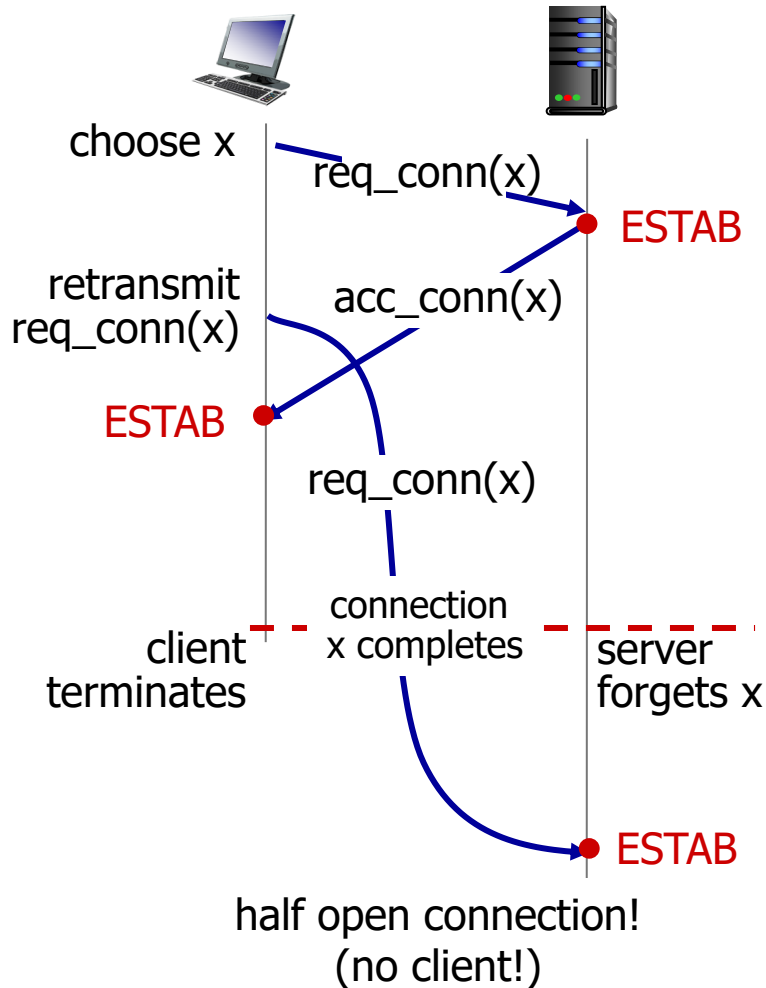
- ❖ variable delays
- ❖ retransmitted messages (e.g. `req_conn(x)`) due to message loss
- ❖ message reordering
- ❖ can't "see" other side

# Problems with Two-way Handshake


- ❑ In an *unreliable* network (e.g. the Internet), lost or delayed segments can cause problems in connection establishment, data transfer and connection termination

# Agreeing to establish a connection

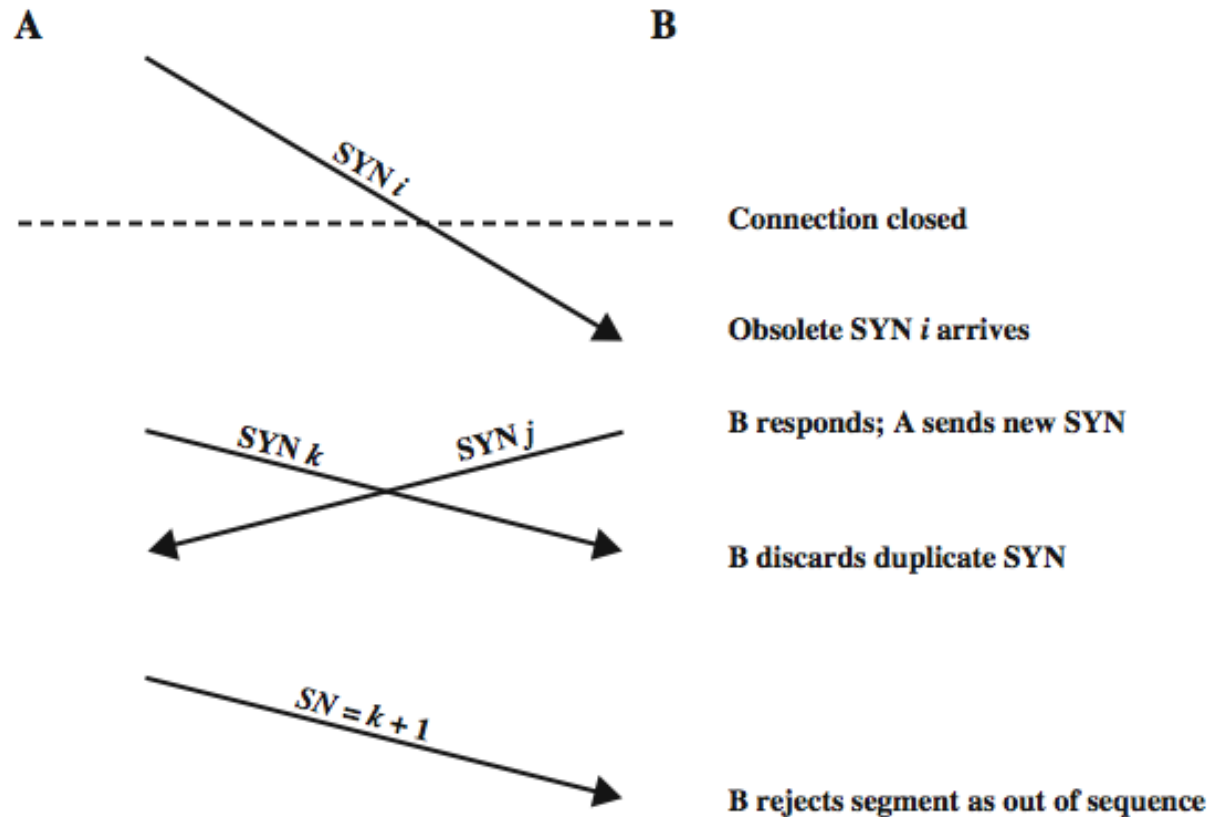
2-way handshake failure scenarios:



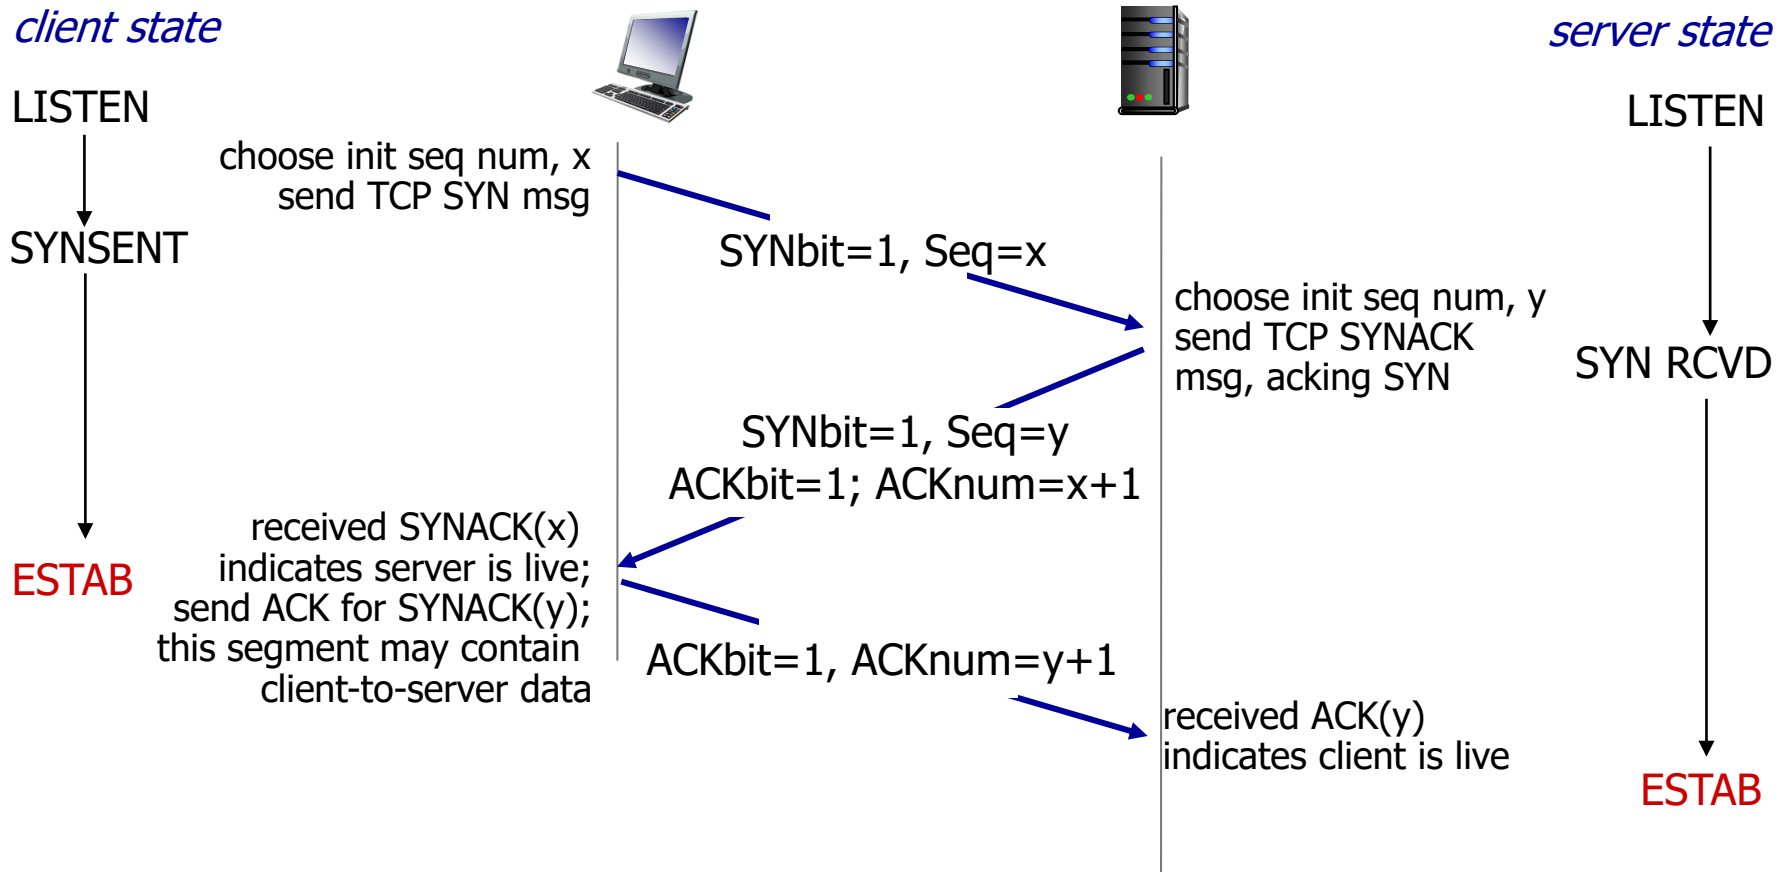
# Two Way Handshake: Obsolete Data Segment



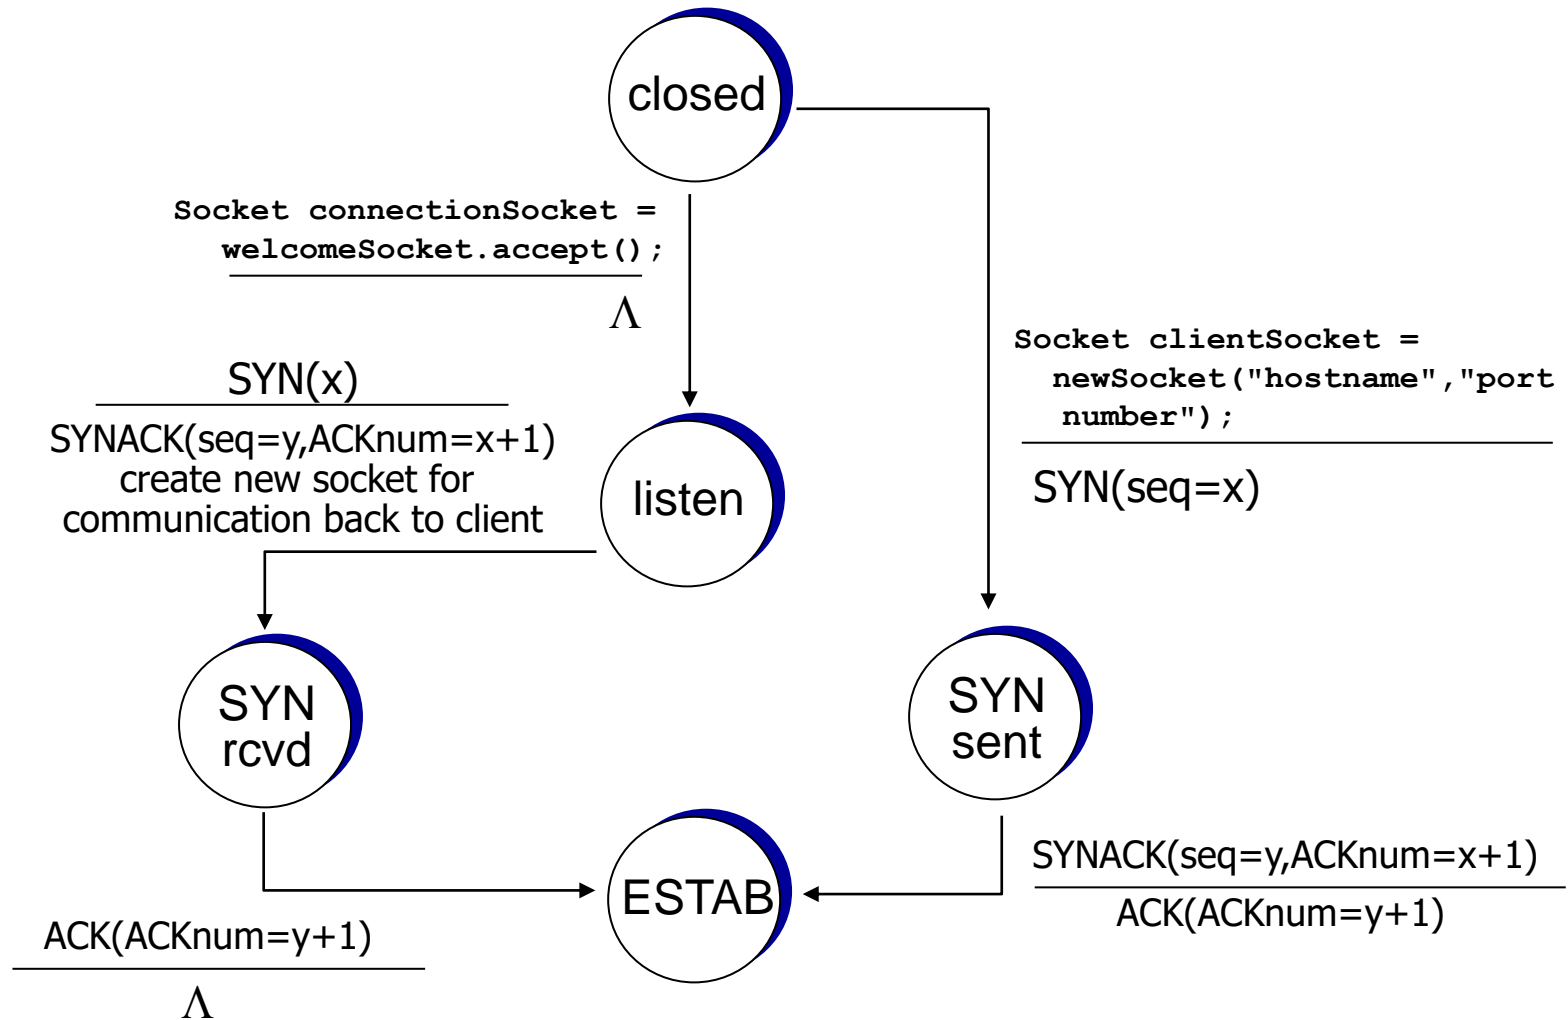
# Two Way Handshake: Obsolete SYN Segment



# TCP 3-way handshake

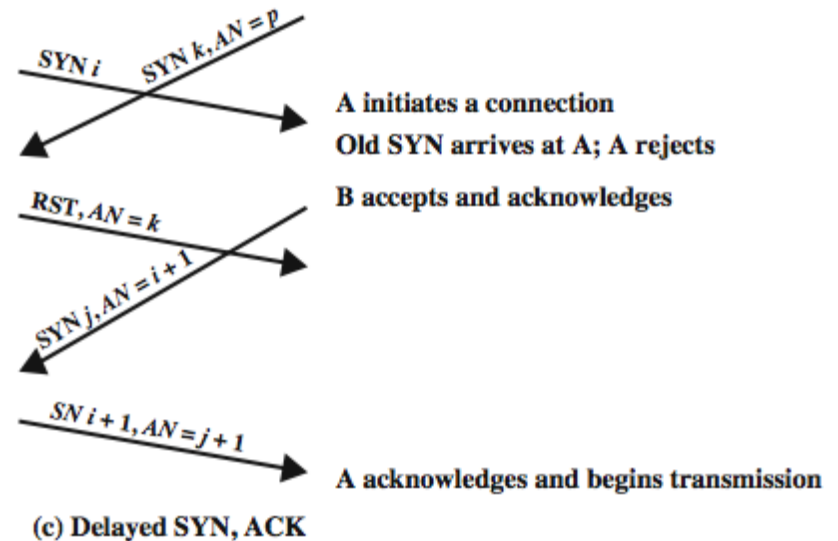




# TCP 3-way handshake: FSM





# Three Way Handshake: Examples



# Closing a connection

## Question:

Do we have a perfect solution for synchronizing the disconnection on both end systems if data can be lost in the network?

See: the two-army problem or  
the Romeo and Juliet problem

## Closing a connection (cont'd)

Answer:

**No, but “three-way handshake” is an acceptable solution.**

# TCP: closing a Connection

- ❑ For better understanding, think of a TCP connection as a pair of simplex connections.
  - "Simplex" means uni-directional data flow
  - Note: A TCP connection is full duplex (i.e., bi-directional.)
- ❑ Each simplex connection is released independently using these two steps:
  - Send a TCP segment with the FIN bit set to one.
  - When the FIN is acknowledged, that direction is shut down.
- ❑ Timers are used for graceful disconnection to avoid the two-army problem.
  - Not a perfect solution, i.e. graceful disconnection cannot be guaranteed
  - In fact, there is no perfect solution at all!

# TCP: closing a connection

*client state*

ESTAB

`clientSocket.close()`

FIN\_WAIT\_1

can no longer  
send but can  
receive data

FIN\_WAIT\_2

wait for server  
close

TIMED\_WAIT

timed wait  
for  $2 * \text{max}$   
segment lifetime

CLOSED



FINbit=1, seq=x

ACKbit=1; ACKnum=x+1

FINbit=1, seq=y

ACKbit=1; ACKnum=y+1

can still  
send data

can no longer  
send data

*server state*

ESTAB

CLOSE\_WAIT

LAST\_ACK

CLOSED

# Summary for connection management

- ❑ Problem: In an *unreliable* network (e.g. the Internet), lost or delayed segments can cause problems in *connection establishment*, *data transfer* and *connection termination*.
- ❑ Acceptable Solution: three way handshake
- ❑ Three way handshake is much better than two way handshake.
- ❑ Timers are used for graceful disconnection to avoid the two-army problem.

# Transport Layer

- ❑ Transport-layer services
- ❑ Connectionless transport: UDP
- ❑ Connection-oriented transport: TCP
  - segment structure
  - reliable data transfer
  - flow control
  - connection management
- ❑ Principles of congestion control
- ❑ TCP congestion control

# Principles of Congestion Control



## Congestion:

- ❑ informally: "too many sources sending too much data too fast for *network* to handle"
- ❑ different from flow control!
- ❑ Signs indicating congestion:
  - lost packets (buffer overflow at routers)
  - long delays (queueing in router buffers)
- ❑ a top-10 problem!




# Causes/costs of congestion: scenario 1

- ❖ two senders, two receivers
- ❖ one router, infinite buffers
- ❖ output link capacity:  $R$
- ❖ no retransmission



- ❖ maximum per-connection throughput:  $R/2$




- ❖ large delays as arrival rate,  $\lambda_{in}$ , approaches capacity

# Causes/costs of congestion: scenario 2



- ❖ four senders/receivers
- ❖ multihop paths
- ❖ timeout/retransmission

Q: what happens as  $\lambda_{in}$  and  $\lambda'_{in}$  increase ?

A: as red  $\lambda'_{in}$  increases, all arriving blue pkts at upper queue are dropped, blue throughput  $\rightarrow 0$



# Causes/costs of congestion: scenario 2



another “cost” of congestion:

- ❖ when a packet dropped, any “upstream transmission capacity” used for that packet was wasted!

# Approaches towards congestion control

Two broad approaches towards congestion control:

## end-end congestion control:

- ❖ no explicit feedback from network
- ❖ congestion inferred from end-system observed loss, delay
- ❖ approach taken by TCP


## network-assisted congestion control:

- ❖ routers provide feedback to end systems
  - single bit indicating congestion (SNA, DECbit, TCP/IP ECN, ATM)
  - explicit rate for sender to send at

# Explicit Congestion Notification (ECN)

## *network-assisted congestion control:*

- two bits in ToS (Type of Service) field of IP header marked *by network router* to indicate congestion
- congestion indication carried to receiving host
- receiver (seeing congestion indication in IP datagram) ) sets ECN-Echo (ECE) bit on receiver-to-sender ACK segment to notify sender of congestion



# Transport Layer

- ❑ Transport-layer services
- ❑ Connectionless transport: UDP
- ❑ Connection-oriented transport: TCP
  - segment structure
  - reliable data transfer
  - connection management
- ❑ Principles of congestion control
- ❑ TCP congestion control


# TCP congestion control:

- ❑ *goal*: TCP sender should transmit as fast as possible, but without congesting network
  - Q: how to find rate *just* below congestion level?
- ❑ decentralized: each TCP sender sets its own rate, based on *implicit* feedback:
  - *ACK*: segment received (a good thing!), network not congested, so increase sending rate
  - *lost segment*: assume loss due to congested network, so decrease sending rate

# TCP congestion control : additive increase multiplicative decrease


- ❖ *approach*: sender increases transmission rate (window size), probing for usable bandwidth, until loss occurs
  - *additive increase*: increase **cwnd** by 1 Maximum Segment Size (MSS) every RTT until loss detected
  - *multiplicative decrease*: cut **cwnd** in half after loss

AIMD saw tooth behavior: probing for bandwidth





# TCP Congestion Control: details



- ❖ sender limits transmission:

$$\text{LastByteSent} - \text{LastByteAcked} \leq \text{cwnd}$$

- ❖ **cwnd** is dynamic, function of perceived network congestion


*TCP sending rate:*

- ❖ *roughly*: send cwnd bytes, wait RTT for ACKS, then send more bytes

$$\text{rate} \approx \frac{\text{cwnd}}{\text{RTT}} \text{ bytes/sec}$$

# TCP Slow Start

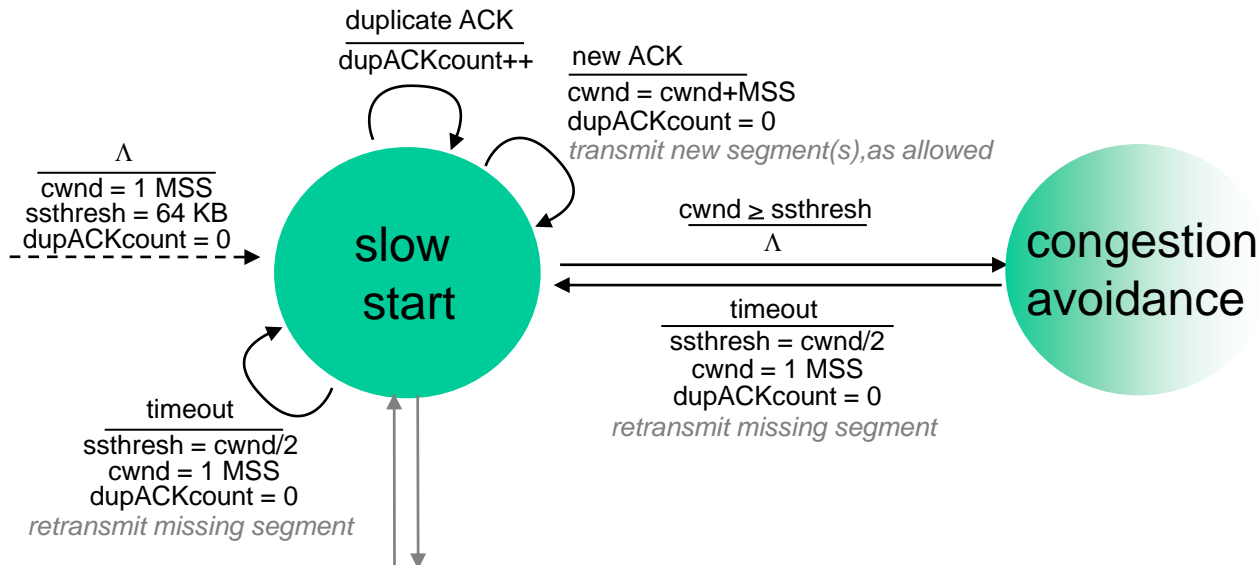
- ❖ when connection begins, increase rate exponentially until first loss event:
  - initially **cwnd** = 1 MSS
  - double **cwnd** every RTT
  - done by incrementing **cwnd** for every ACK received
- ❖ summary: initial rate is slow but ramps up exponentially fast



# Transitioning into/out of slowstart


**ssthresh**: cwnd threshold maintained by TCP

- ❑ on loss event: set ssthresh to cwnd/2
  - remember (half of) TCP rate when congestion last occurred
- ❑ when  $\text{cwnd} \geq \text{ssthresh}$ : transition from slowstart to congestion avoidance phase




# TCP: congestion avoidance

- when `cwnd`  $\geq$  `ssthresh`, `cwnd` grows linearly
  - increase `cwnd` by 1 MSS per RTT
  - approach possible congestion slower than in slowstart



(a) Slow start, ending with a timeout



(b) Slow start followed by congestion avoidance


# TCP Congestion Control: segment loss event

- ❑ Loss indicated by timeout:
  - cut `cwnd` to 1 MSS
  - `cwnd` then grows exponentially (as in slow start) to threshold, then grows linearly
- ❑ Loss indicated by 3 duplicate ACKs: TCP RENO
  - `cwnd` is cut in half and `cwnd` then grows linearly: less aggressively than on timeout (**Fast Recovery**)
  - Note that TCP Tahoe always set `cwnd` to 1 (timeout or 3 duplicate acks)

## Philosophy:

- ❑ 3 dup ACKs indicates network capable of delivering some segments (recall **fast retransmit**)
- ❑ timeout indicates a "more alarming/serious" congestion scenario

# TCP fast retransmit



# Why Fast Recovery is used?

- ❑ When TCP retransmits a segment using **Fast Retransmit**, a segment was assumed lost
- ❑ Some congestion avoidance measures are appropriate at this point
- ❑ **Slow Start** may be unnecessarily conservative since multiple acks indicate segments are getting through (meaning congestion not so serious)
- ❑ **Fast Recovery**: retransmit lost segment, cut CongWin in half, and proceed with linear increase of CongWin (avoiding “slow” start-up)



# TCP Congestion Control: ACK received

## ACK received: increase CWND

- ❑ slowstart phase:
  - increase exponentially fast (despite name) at connection start or following timeout
- ❑ congestion avoidance:
  - increase linearly

# Variants of TCP Congestion Control Schemes

- ❑ TCP Tahoe: Slow Start + Congestion Avoidance.
- ❑ TCP Reno: TCP Tahoe + Fast Retransmit + Fast Recovery.

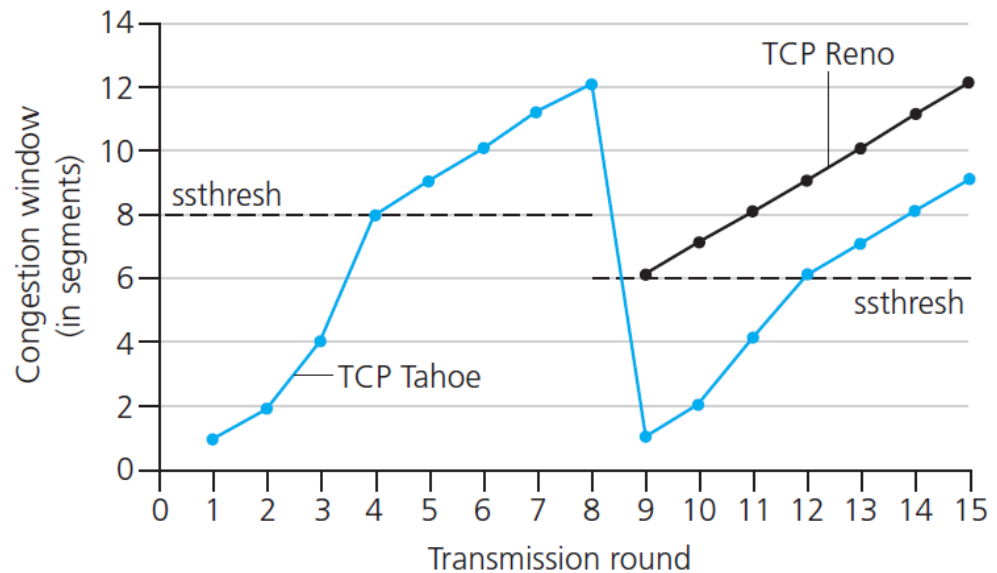
# TCP: from slow start to congestion avoidance

**Q:** when should the exponential increase switch to linear?

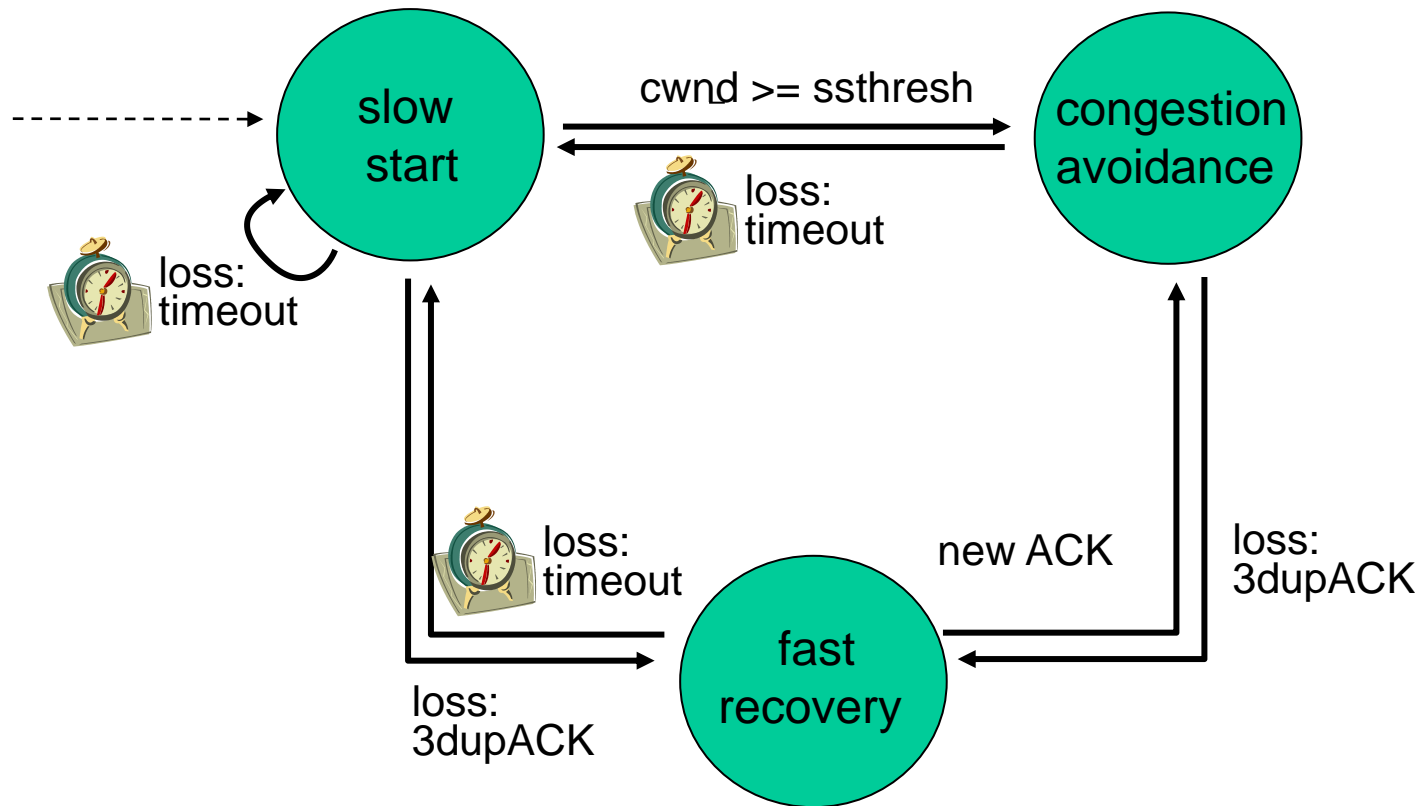
**A:** when **cwnd** gets to 1/2 of its value before last timeout.

## Implementation:

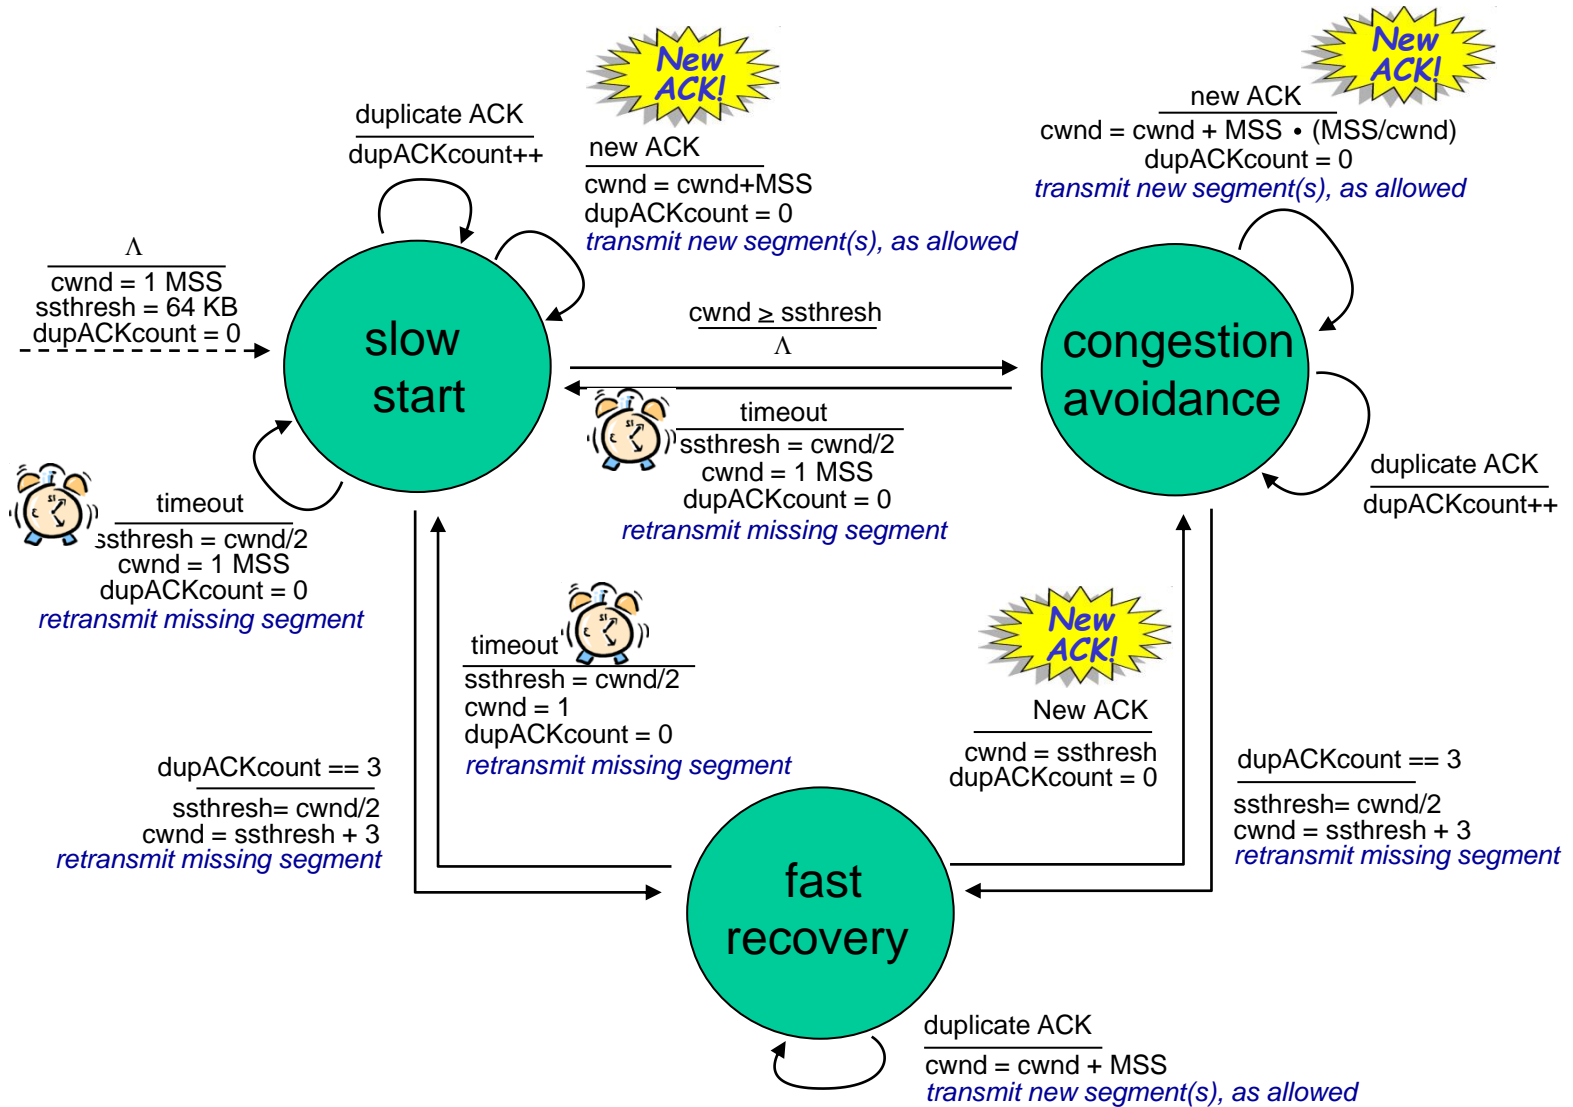
- ❖ variable **ssthresh**
- ❖ on loss event, **ssthresh** is set to 1/2 of **cwnd** just before loss event



# TCP Reno congestion control FSM: overview



# TCP Reno congestion control FSM: details



# Summary: TCP Congestion Control


- ❑ when  $cwnd < ssthresh$ , sender is in **slow-start** phase, window grows exponentially.
- ❑ when  $cwnd \geq ssthresh$ , sender is in **congestion-avoidance** phase, window grows linearly.
- ❑ when **triple duplicate ACK** occurs,  $ssthresh$  set to  $cwnd/2$ ,  $cwnd$  set to  $ssthresh$  and go to **congestion-avoidance** phase.
- ❑ when **timeout** occurs,  $ssthresh$  set to  $cwnd/2$ ,  $cwnd$  set to 1 MSS and go to **slow-start** phase.

# TCP sender congestion control

State	Event	TCP Sender Action	Commentary
Slow Start (SS)	ACK receipt for previously unacked data	$\text{CongWin} = \text{CongWin} + 1 \text{ MSS}$ , If $(\text{CongWin} \geq \text{Threshold})$ set state to "Congestion Avoidance"	Resulting in a doubling of CongWin every RTT
Congestion Avoidance (CA)	ACK receipt for previously unacked data	$\text{CongWin} = \text{CongWin} + 1 \text{ MSS} * (\text{MSS}/\text{CongWin})$	Additive increase, resulting in increase of CongWin by 1 MSS every RTT
SS or CA	Loss event detected by triple duplicate ACK	$\text{Threshold} = \text{CongWin}/2$ , $\text{CongWin} = \text{Threshold}$ , Set state to "Congestion Avoidance"	Fast recovery, implementing multiplicative decrease. CongWin will not drop below 1 MSS.
SS or CA	Timeout	$\text{Threshold} = \text{CongWin}/2$ , $\text{CongWin} = 1 \text{ MSS}$ , Set state to "Slow Start"	Enter slow start
SS or CA	Duplicate ACK	Increment duplicate ACK count for segment being acked	CongWin and Threshold not changed

# TCP Fairness

**fairness goal:** if  $K$  TCP sessions share same bottleneck link of bandwidth  $R$ , each should have average rate of  $R/K$






# Why is TCP fair?

two competing sessions:

- ❖ additive increase gives slope of 1, as throughput increases
- ❖ multiplicative decrease decreases throughput proportionally



# Fairness (more)

## Fairness and UDP

- ❑ multimedia apps often do not use TCP
  - do not want rate slowed down by congestion control
- ❑ instead use UDP:
  - pump audio/video at constant rate, tolerate packet loss

## Fairness and parallel TCP connections

- ❑ nothing prevents app from opening parallel connections between 2 hosts.
- ❑ web browsers do this
- ❑ For example: link of rate  $R$  supporting 9 connections;
  - new app asks for 1 TCP, gets rate  $R/10$
  - new app asks for 9 TCPs, gets rate  $R/2$  !

# Summary

- ❑ Transport-layer services
- ❑ Connectionless transport: UDP
- ❑ Connection-oriented transport: TCP
  - segment structure
  - reliable data transfer
  - flow control
  - connection management
- ❑ Principles of congestion control
- ❑ TCP congestion control

Q & A