# SDSC3002 Intro to Data Mining

Yu Yang
yuyang@city.edu.hk

# Outline

What is Data Mining?

SDSC3002

# Data, Information & Knowledge

- Data
  - Phenomenon observed by people, due to some underlying mechanism
- Information
  - Organized data that has meaning and value
- Knowledge
  - The concept of understanding information based on identified patterns that provide insights to applications

# Data Mining

- **K**nowledge **D**iscovery from **D**ata
  - Synonym of data mining
  - ACM KDD is the best data mining conference
- Identifying interesting patterns/knowledge from (big) data
- Example: Recommender Systems
  - Data: users' browsing/purchasing/like data
  - Knowledge: a mapping function that maps a (user, item) pair to a clicking/purchasing probability
  - Insight: displaying related and personalized items to users to enhance user stickiness and increase revenue

# Demand for Data Mining

**Demand for deep analytical talent in the United States could be 50 to 60 percent greater than its projected supply by 2018**

Supply and demand of deep analytical talent by 2018
Thousand people

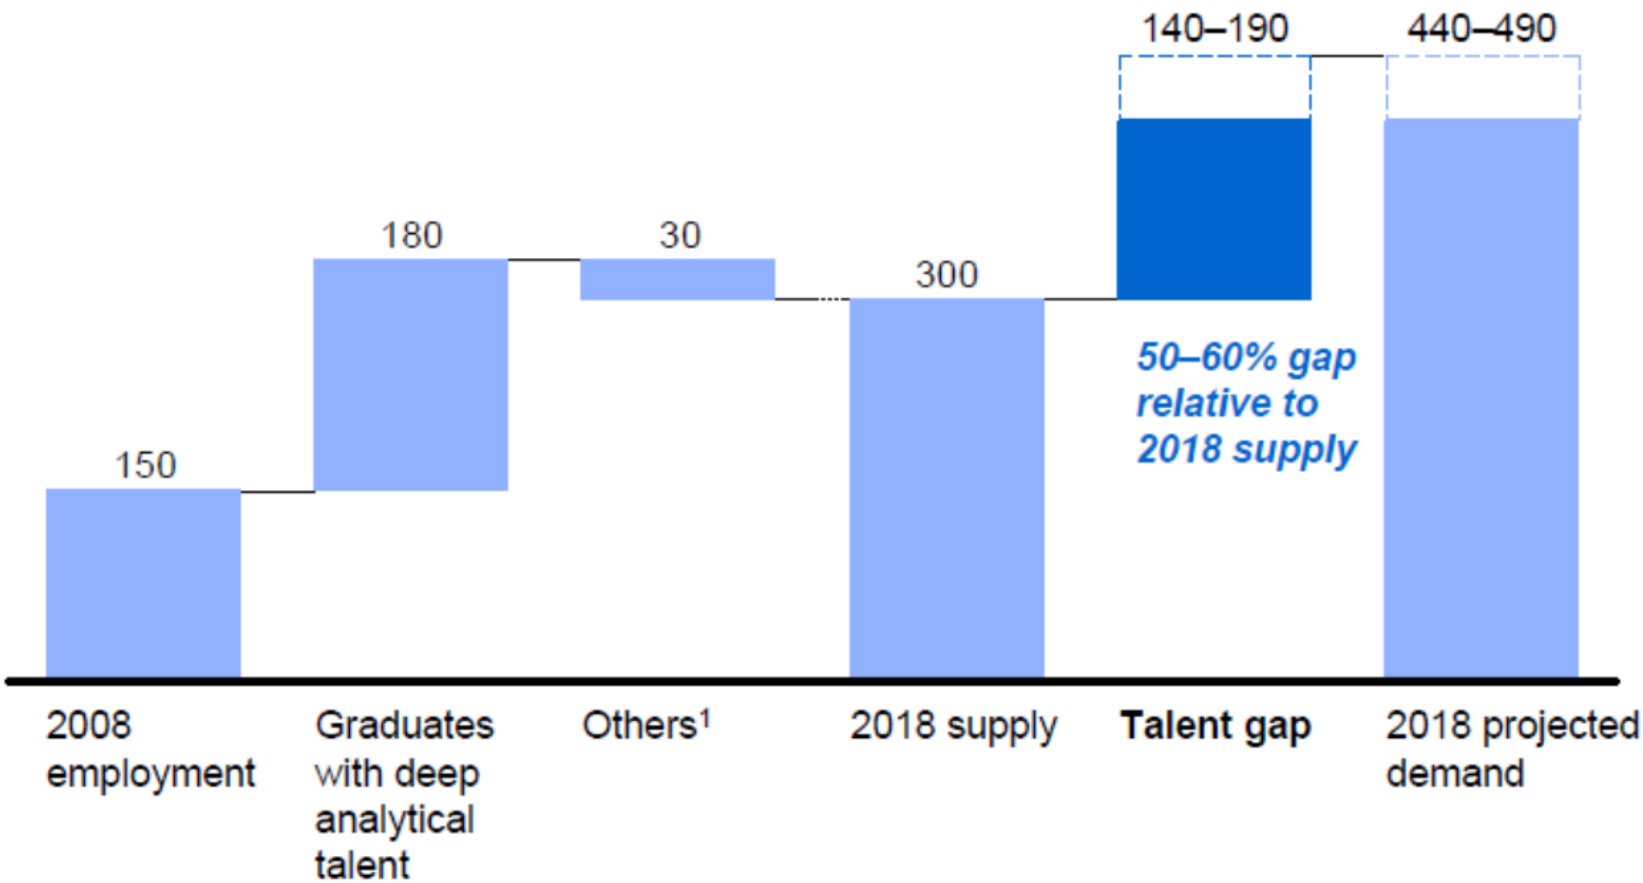

# Data Mining Tasks

- Descriptive task
  - Find human-interpretable patterns that describe the data
  - Examples: itemset mining, community detection
- Predictive task
  - Use some variables to predict unknown or future values of other variables
  - Examples: spam email detection, recommender systems
- Complex task
  - Combine descriptive/predictive methods with decision making methods
  - Examples: item bundling, warehouse inventory design

# Example: Data-Driven Distribution Network Design

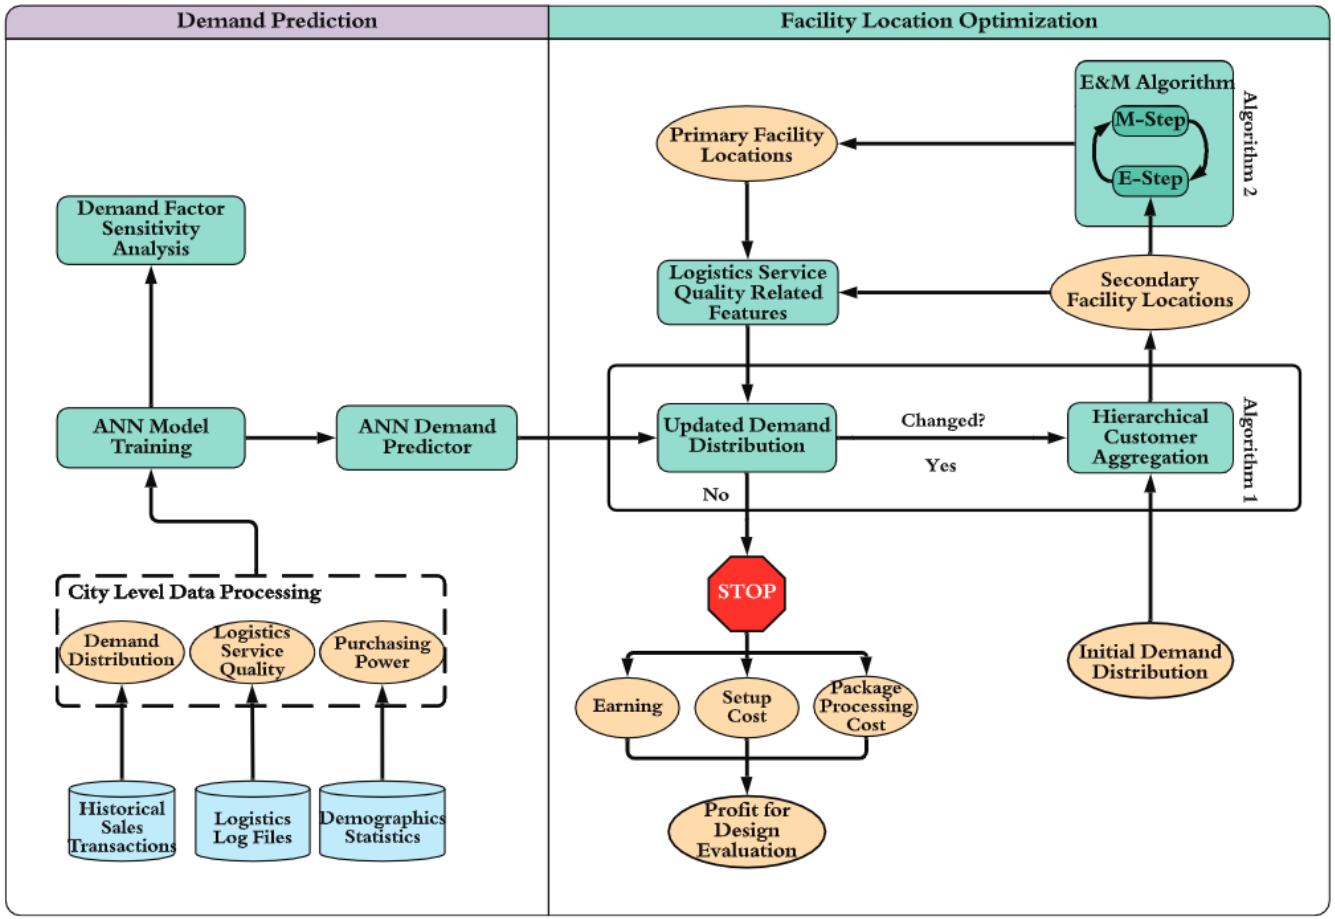


Liu, Junming, et al. "Iterative Prediction-and-Optimization for E-Logistics Distribution Network Design." INFORMS Journal on Computing (2021).

# Data Mining: Cultures

- Sub-communities in Data Mining
  - Database: big data, "simple" queries
    - DM as complex queries, scalability/efficiency
  - Machine Learning: small data, complex models
    - DM as inference of models, prediction accuracy
  - Theory: (randomized) algorithms

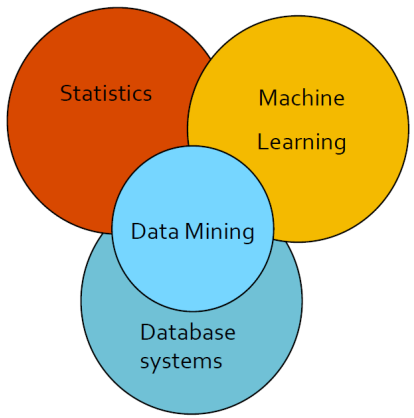

# Related Courses

- SDSC2102 Statistical Methods and Data Analysis
  - 3002 covers non-statistical methods and more "practical" methods
- SDSC3006 Fundamentals of Machine Learning I
  - Machine learning is a very important tool in Data Mining, large overlap
  - "practical" machine learning
  - More business applications
- SDSC3001 Big Data: The Arts and Science of Scaling
  - 3002 is more application-driven

# Outline

What is Data Mining?

SDSC3002

# Contents


- Data Mining Methodologies
  - Similarity/Distance
  - Clustering
  - Graph Mining
  - Data Privacy
- Data Mining Applications
  - Market Basket Analysis
  - Recommender Systems
  - Online Advertising
  - Social Network

# Course Logistics

- Course Website: `https://canvas.cityu.edu.hk/courses/46765`
- Office hour: 14:50-15:50 every Wednesday (starting from week 2)
- Textbook: Mining of Massive Datasets, *by Jure Leskovec, Anand Rajaraman, Jeff Ullman*
- Grading Scheme
  - 4 Assignments: 10% each
  - Group Project (max. 4 students per group): 20%
  - Midterm: 10%
  - Final exam: 30%

# Acknowledgement


- Some of the contents originate from Jure Leskovec's slides for CS246Stanford