Lab 1:
Introduction to R

Outline

- Install R and data packages
- Basic commands
- Example of preliminary analysis of a dataset

Install R and data packages

Install R

- https://www.r-project.org/

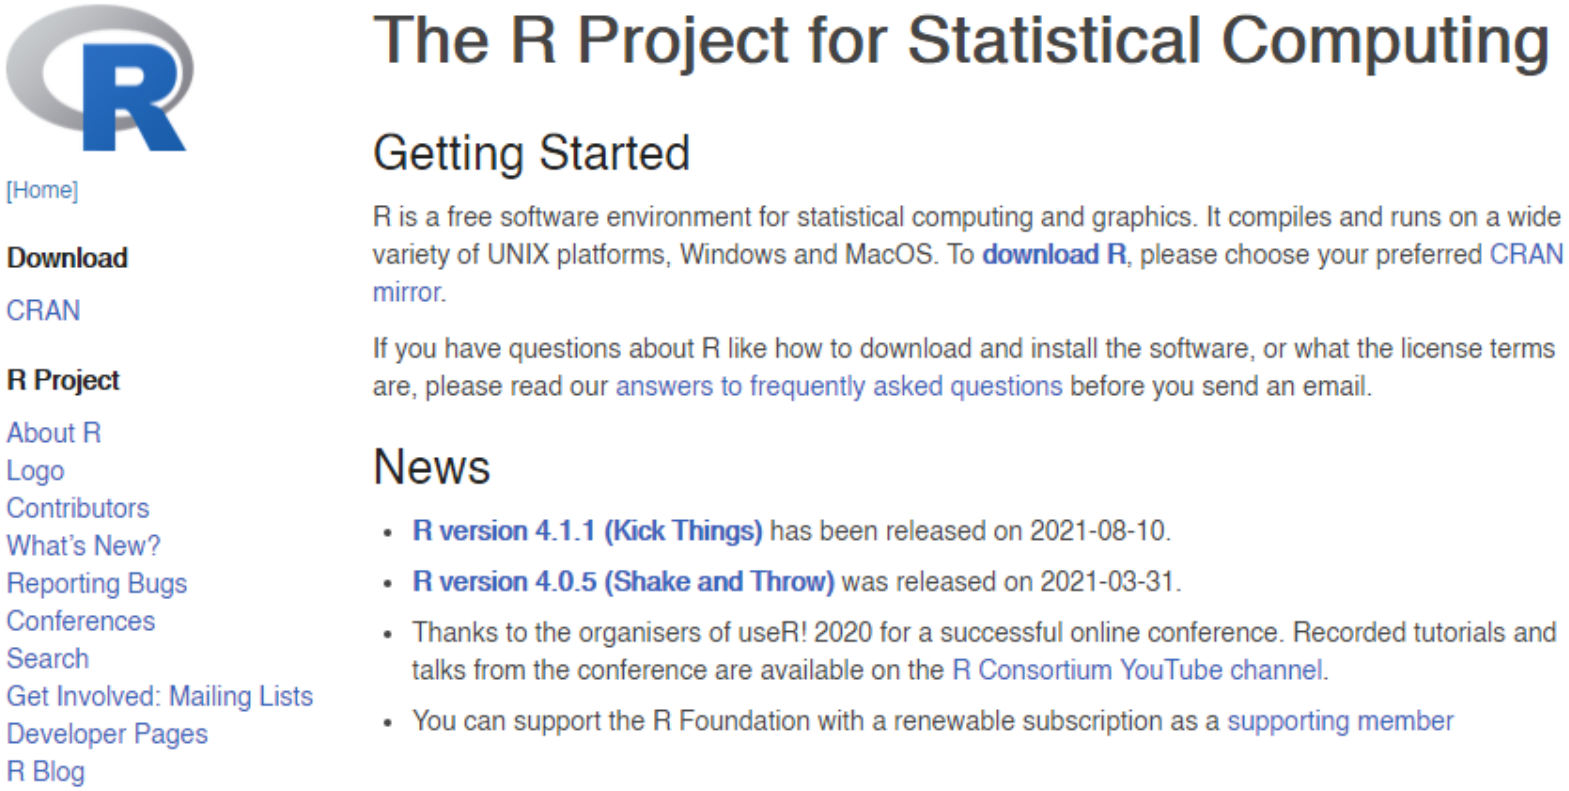

[Home]

Download

CRAN

R Project

About R
Logo
Contributors
What's New?
Reporting Bugs
Conferences
Search
Get Involved: Mailing Lists
Developer Pages
R Blog

The R Project for Statistical Computing

Getting Started

R is a free software environment for statistical computing and graphics. It compiles and runs on a wide variety of UNIX platforms, Windows and MacOS. To **download R**, please choose your preferred CRAN mirror.

If you have questions about R like how to download and install the software, or what the license terms are, please read our answers to frequently asked questions before you send an email.

News

- **R version 4.1.1 (Kick Things)** has been released on 2021-08-10.
- **R version 4.0.5 (Shake and Throw)** was released on 2021-03-31.
- Thanks to the organisers of useR! 2020 for a successful online conference. Recorded tutorials and talks from the conference are available on the R Consortium YouTube channel.
- You can support the R Foundation with a renewable subscription as a supporting member

Install Rstudio Desktop （Optional）

- https://www.rstudio.com/products/rstudio/download/

RStudio Desktop 1.4.1717 - Release Notes

1. Install R. RStudio requires R 3.0.1+.

2. Download RStudio Desktop. Recommended for your system:

DOWNLOAD RSTUDIO FOR WINDOWS
1.4.1717 | 156.18MB

Requires Windows 10 (64-bit)

Data Sets Used in Labs and Exercises

- ISLR2:
  - Data for an Introduction to Statistical Learning with Applications in R
  - https://cran.r-project.org/package=ISLR2
- The MASS library
  - Functions and datasets to support Venables and Ripley's MASS
  - https://cran.r-project.org/package=MASS
- Base R

Name	Description
Auto	Gas mileage, horsepower, and other information for cars.
Bikeshare	Hourly usage of a bike sharing program in Washington, DC.
Boston	Housing values and other information about Boston census tracts.
BrainCancer	Survival times for patients diagnosed with brain cancer.
Caravan	Information about individuals offered caravan insurance.
Carseats	Information about car seat sales in 400 stores.
College	Demographic characteristics, tuition, and more for USA colleges.
Credit	Information about credit card debt for 10,000 customers.
Default	Customer default records for a credit card company.
Fund	Returns of 2,000 hedge fund managers over 50 months.
Hitters	Records and salaries for baseball players.
Khan	Gene expression measurements for four cancer types.
NCI60	Gene expression measurements for 64 cancer cell lines.
NYSE	Returns, volatility, and volume for the New York Stock Exchange.
OJ	Sales information for Citrus Hill and Minute Maid orange juice.
Portfolio	Past values of financial assets, for use in portfolio allocation.
Publication	Time to publication for 244 clinical trials.
Smarket	Daily percentage returns for S&P 500 over a 5-year period.
USArrests	Crime statistics per 100,000 residents in 50 states of USA.
Wage	Income survey data for men in central Atlantic region of USA.
Weekly	1,089 weekly stock market returns for 21 years.

Install ISLR2 Package

- Manual download with R
  - Click Packages & Data → Package Installer → input the package name → select a mirror site
- Manual download with Rstudio
  - Click Tools → Install package → install from CRAN/local archive file
- By R command line
  - install.packages("ISLR2")

Basic commands

Vector

- Insert vector using function c()
- Check length of vector using length()

```
> x <- c(1, 2, 3, 4, 5, 6, 7, 8, 9)
> x
[1] 1 2 3 4 5 6 7 8 9
> length(x)
[1] 9
```

Matrix

- Declare a matrix using function matrix()
- Use byrow =TRUE/FALSE to specify order
- Use dim() to find dimension of a matrix

Matrix

```
> x <- matrix (data = c(1, 2, 3, 4, 5, 6), nrow = 2, ncol = 3)
> x
     [,1] [,2] [,3]
[1,]    1    3    5
[2,]    2    4    6
> x <- matrix (data = c(1, 2, 3, 4, 5, 6), nrow = 2, ncol = 3, byrow = TRUE)
> x
     [,1] [,2] [,3]
[1,]    1    2    3
[2,]    4    5    6
> x <- matrix (data = c(1, 2, 3, 4, 5, 6), nrow = 2, ncol = 3, byrow = FALSE)
> x
     [,1] [,2] [,3]
[1,]    1    3    5
[2,]    2    4    6
> dim(x)
[1] 2 3
```

Select Elements in A Matrix

```
> A <- matrix(1:16, 4, 4)
> A
     [,1] [,2] [,3] [,4]
[1,]    1    5    9   13
[2,]    2    6   10   14
[3,]    3    7   11   15
[4,]    4    8   12   16
> A[2, 3]
[1] 10
> A[c(1, 3), c(2, 4)]
     [,1] [,2]
[1,]    5   13
[2,]    7   15
> A[1:2,]
     [,1] [,2] [,3] [,4]
[1,]    1    5    9   13
[2,]    2    6   10   14
```

```
> A[, 1]
[1] 1 2 3 4
> A[, 1:2]
     [,1] [,2]
[1,]    1    5
[2,]    2    6
[3,]    3    7
[4,]    4    8
> A[-1,]
     [,1] [,2] [,3] [,4]
[1,]    2    6   10   14
[2,]    3    7   11   15
[3,]    4    8   12   16
> A[-c(1, 2),]
     [,1] [,2] [,3] [,4]
[1,]    3    7   11   15
[2,]    4    8   12   16
```

Generate Random Numbers

- Generate random numbers from a standard normaldistribution using rnorm (n)

```
> y <- rnorm(20)
> y
 [1] -1.09460899  0.22386861 -0.20583813 -1.11530919  0.58994271 -2.06441523  1.39271334
 [8] -0.23312401 -0.10541311 -0.63185659  0.25970922  0.43194340 -0.04194421  0.09849715
[15] -0.50593705  1.15531491 -1.14990503 -0.26525399  0.87302274  0.12407061
```

- Calculate mean(), var(), sd() of random numbers

```
> mean(y)
[1] -0.1132261
> var(y)
[1] 0.7009881
> sd(y)
[1] 0.8372503
```

Set the Seed of Random Number Generator

- Set the seed of random number generator using set.seed
- To reproduce the exact same set of random numbers, use the same seed

```
> set.seed(1)
> rnorm(5)
[1] -0.6264538  0.1836433 -0.8356286  1.5952808  0.3295078
> rnorm(5)
[1] -0.8204684  0.4874291  0.7383247  0.5757814 -0.3053884
> set.seed(1)
> rnorm(5)
[1] -0.6264538  0.1836433 -0.8356286  1.5952808  0.3295078
```

Example of preliminary analysis of a dataset

Load Dataset

- To load a data set in the ISLR2 package or other packages/libraries, you only need to load the package

```
> library (ISLR2)
```

- To load an external data set, first specify the directory, Misc → Change Working Directory
  - If the data are saved as a text file

```
> Auto <- read.table("Auto.data", header=T, na.strings="?", stringsAsFactors=T)
```

  - If the data are saved as a csv file (Excel)

```
> College <- read.csv("College.csv", na.strings="?", stringsAsFactors=T)
```

- Try loading external data files using datasets available on the textbook website https://www.statlearning.com/resources-second-edition

View Data

Then, check a dataset by typing its name in console or view()

	mpg	cylinders	displacement	horsepower	weight	acceleration	year	origin	name
1	18	8	307.0	130	3504	12.0	70	1	chevrolet chevelle malibu
2	15	8	350.0	165	3693	11.5	70	1	buick skylark 320
3	18	8	318.0	150	3436	11.0	70	1	plymouth satellite
4	16	8	304.0	150	3433	12.0	70	1	amc rebel sst
5	17	8	302.0	140	3449	10.5	70	1	ford torino
6	15	8	429.0	198	4341	10.0	70	1	ford galaxie 500
7	14	8	454.0	220	4354	9.0	70	1	chevrolet impala
8	14	8	440.0	215	4312	8.5	70	1	plymouth fury iii
9	14	8	455.0	225	4425	10.0	70	1	pontiac catalina
10	15	8	390.0	190	3850	8.5	70	1	amc ambassador dpl
11	15	8	383.0	170	3563	10.0	70	1	dodge challenger se
12	14	8	340.0	160	3609	8.0	70	1	plymouth 'cuda 340

Dimension & Variables

- Type names(datasetname), e.g., names(Auto), to list all attributes (column names) of the table

```
> names(Auto)
[1] "mpg"          "cylinders"    "displacement" "horsepower"
[5] "weight"       "acceleration" "year"         "origin"
[9] "name"
> dim(Auto)
[1] 397   9
> Auto <- na.omit(Auto)
> dim(Auto)
[1] 392   9
```

Background

- To gather more information about the data set, type ?datasetname (e.g., ?Auto, or help(Auto))

Auto Data Set

Description

Gas mileage, horsepower, and other information for 397 vehicles.

Usage

```
Auto
```

Format

A data frame with 397 observations on the following 9 variables.

`mpg`

miles per gallon

`cylinders`

Number of cylinders between 4 and 8

`displacement`

Engine displacement (cu. inches)

`horsepower`

Engine horsepower

Access Variables in the Dataset

- Method 1: giving name of the variable and the dataset

```
> Auto$mpg
  [1] 18.0 15.0 18.0 16.0 17.0 15.0 14.0
  [8] 14.0 14.0 15.0 15.0 14.0 15.0 14.0
 [15] 24.0 22.0 18.0 21.0 27.0 26.0 25.0
```

- Method 2: first attach the dataset to the R search path, then all variables in the dataset can be accessed by simply giving their names

```
> attach(Auto)
> mpg
  [1] 18.0 15.0 18.0 16.0 17.0 15.0 14.0 14.0
  [9] 14.0 15.0 15.0 14.0 15.0 14.0 24.0 22.0
 [17] 18.0 21.0 27.0 26.0 25.0 24.0 25.0 26.0
```

Numerical Summaries

- Numerical summary of dataset or variable summary(dataset) or summary(dataset$colname)

```
> summary(Auto)
      mpg          cylinders      displacement     horsepower        weight
 Min.   : 9.00   Min.   :3.000   Min.   : 68.0   Min.   : 46.0   Min.   :1613
 1st Qu.:17.00   1st Qu.:4.000   1st Qu.:105.0   1st Qu.: 75.0   1st Qu.:2225
 Median :22.75   Median :4.000   Median :151.0   Median : 93.5   Median :2804
 Mean   :23.45   Mean   :5.472   Mean   :194.4   Mean   :104.5   Mean   :2978
 3rd Qu.:29.00   3rd Qu.:8.000   3rd Qu.:275.8   3rd Qu.:126.0   3rd Qu.:3615
 Max.   :46.60   Max.   :8.000   Max.   :455.0   Max.   :230.0   Max.   :5140

  acceleration        year           origin                      name
 Min.   : 8.00   Min.   :70.00   Min.   :1.000   amc matador       :  5
 1st Qu.:13.78   1st Qu.:73.00   1st Qu.:1.000   ford pinto        :  5
 Median :15.50   Median :76.00   Median :1.000   toyota corolla    :  5
 Mean   :15.54   Mean   :75.98   Mean   :1.577   amc gremlin       :  4
 3rd Qu.:17.02   3rd Qu.:79.00   3rd Qu.:2.000   amc hornet        :  4
 Max.   :24.80   Max.   :82.00   Max.   :3.000   chevrolet chevette:  4
                                                 (Other)           :365
```

Graphs

- Generate figures of single variables using plot()

```
> plot(College$F.Undergrad, type='l', xlab="College Index", ylab="#Full-time Undergraduates",
main="Figure 1")
```

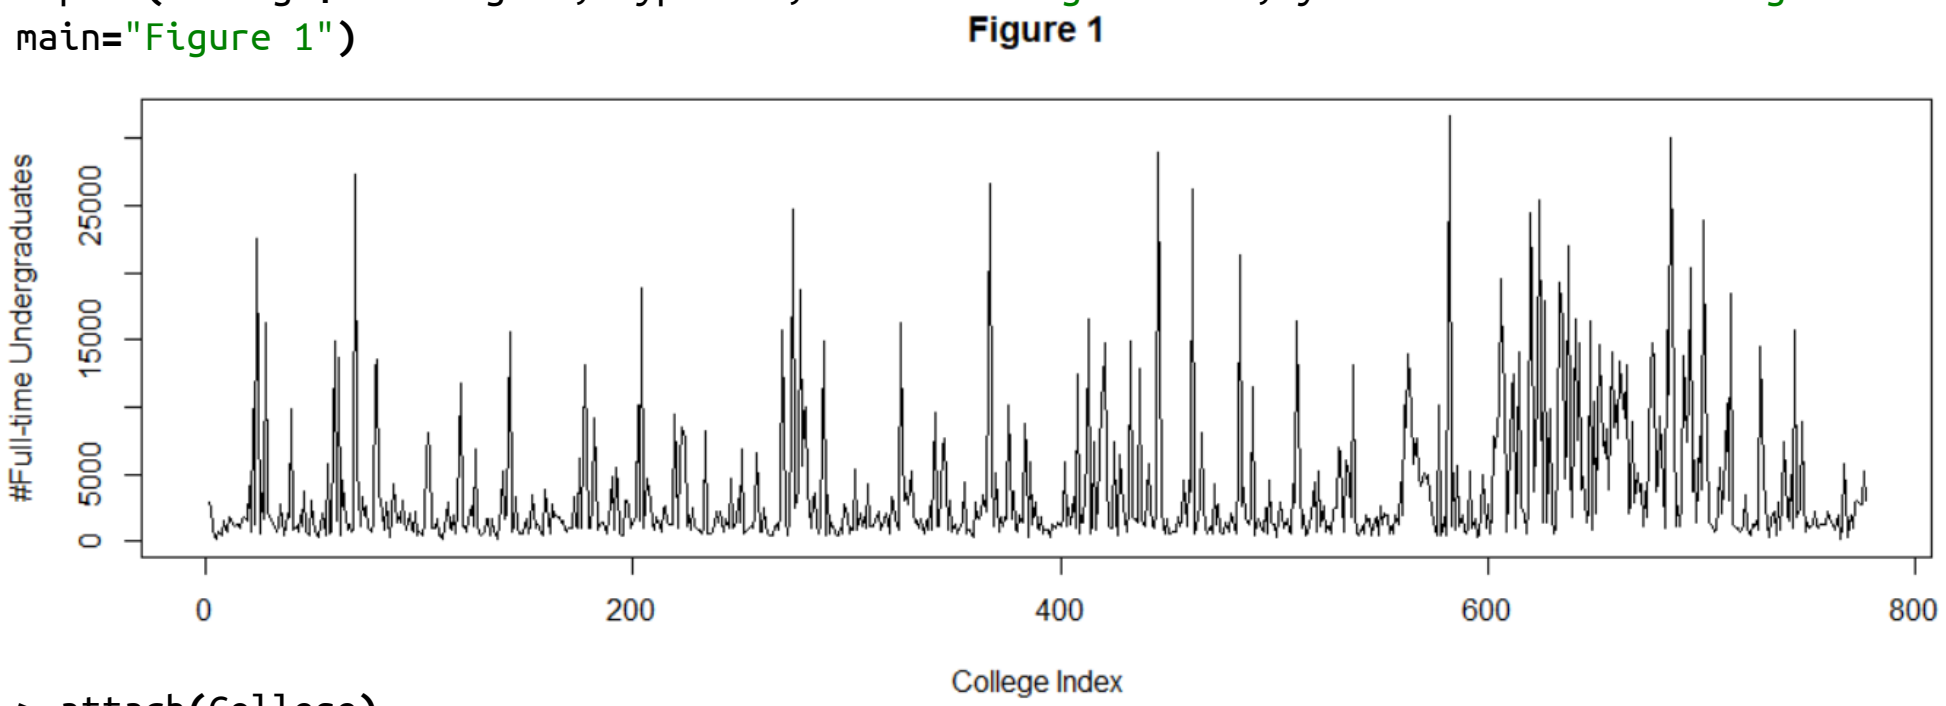


```
> attach(College)
> plot(F.Undergrad, type='l', xlab="College Index", ylab="#Full-time Undergraduates", main="Figure 1")
```

Graphs

- Generate plots of two variables

```
> plot(Apps, Accept, xlab="#Received",
ylab="#Accepted")
```

`Private` : Public/private indicator

`Apps` : Number of applications received

`Accept` : Number of applicants accepted

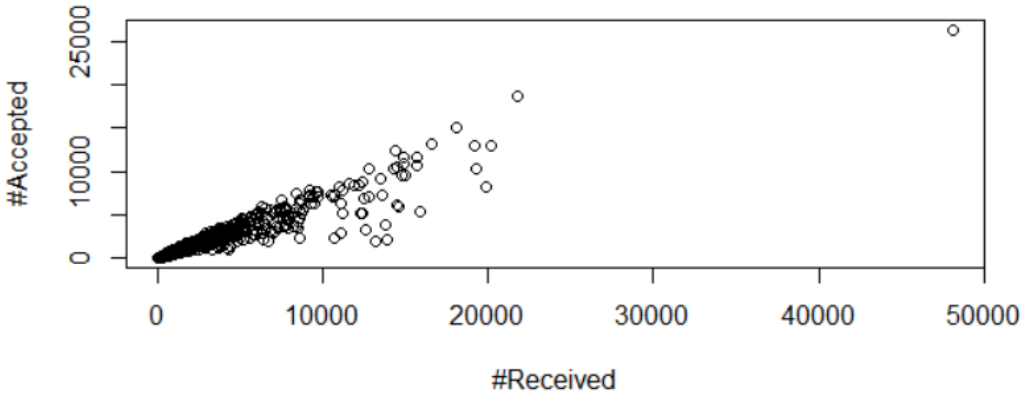


```
> plot(Apps[Private=="Yes"], Accept[Private=="Yes"], col="blue", xlab="#Received", ylab="#Accepted")
> points(Apps[Private=="No"], Accept[Private=="No"], col="green", xlab="#Received", ylab="#Accepted")
```