Topic 9. Deep Learning

➢ Neural network (NN) appeared in late 1980s.

➢ NN resurfaced around 2010 with the new name deep learning.

"Review of Machine Learning and Deep Learning Application in Mine Microseismic Event Classification", *Mining of Mineral Deposits*, 15:19-26.
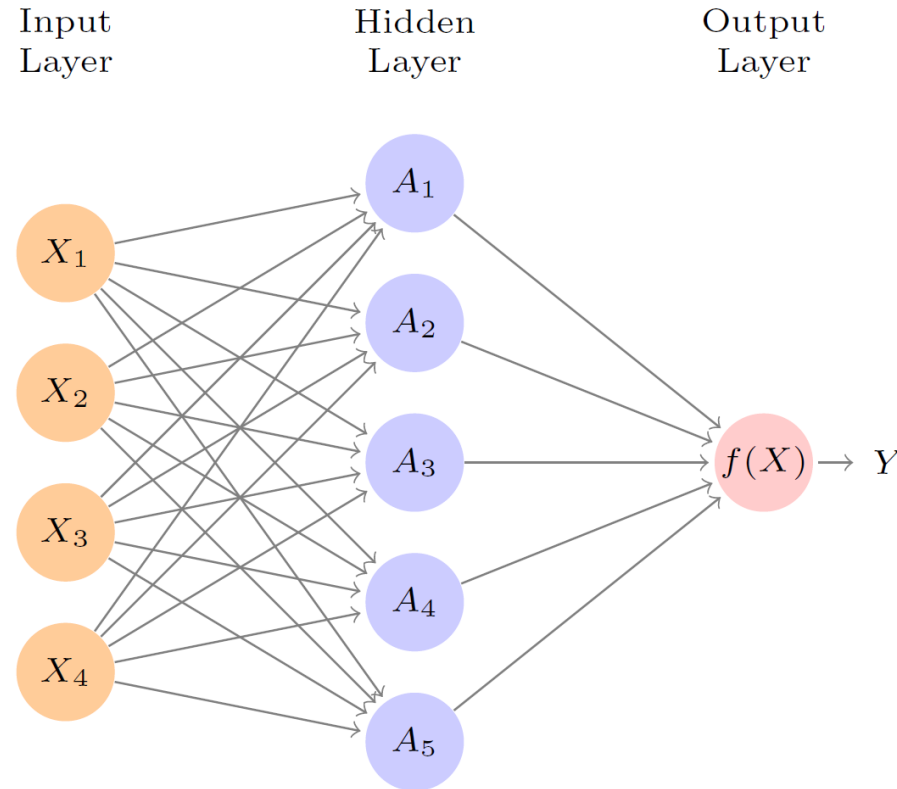
Outline

➤ Basics: single layer neural network and fitting

➤ More complex neural networks

➤ When to use deep learning

Basics:
Single Neural Network and Fitting

Single Layer Neural Networks

Input
Layer

Hidden
Layer

Output
Layer

X_1

X_2

X_3

X_4

A_1

A_2

A_3

A_4

A_5

$f(X) \rightarrow Y$

$$
\begin{aligned}
f(X) &= \beta_0 + \sum_{k=1}^{K} \beta_k h_k(X) \\
&= \beta_0 + \sum_{k=1}^{K} \beta_k g(w_{k0} + \sum_{j=1}^{p} w_{kj} X_j)
\end{aligned}
$$

Hidden Layer

$$A_k = h_k(X) = g(w_{k0} + \sum_{j=1}^{p} w_{kj} X_j)$$

➢ A_k is a transformation of the original features (predictors)

➢ $g(z)$ is a nonlinear activation function specified in advance

  ➢ Sigmoid activation function

  $$g(z) = \frac{e^z}{1 + e^z} = \frac{1}{1 + e^{-z}}$$

  ➢ ReLU (rectified linear unit) activation function

  $$g(z) = (z)_+ = \begin{cases} 0 & \text{if } z < 0 \\ z & \text{otherwise.} \end{cases}$$

➤ $f(X)$ is a linear function of activations

$$f(X) = \beta_0 + \sum_{k=1}^{K} \beta_k A_k$$

Neurons in Human Brain

➤ The name neural network originally derived from thinking of hidden units as analogous to neurons in brain.

➤ Activations A_k close to 1 are *firing* and close to 0 are *silent*

(dana.org)

➢ A nonlinear least squares problem

$$\operatorname*{minimize}_{\{w_k\}_1^K,\ \beta} \frac{1}{2} \sum_{i=1}^{n} (y_i - f(x_i))^2,$$

where

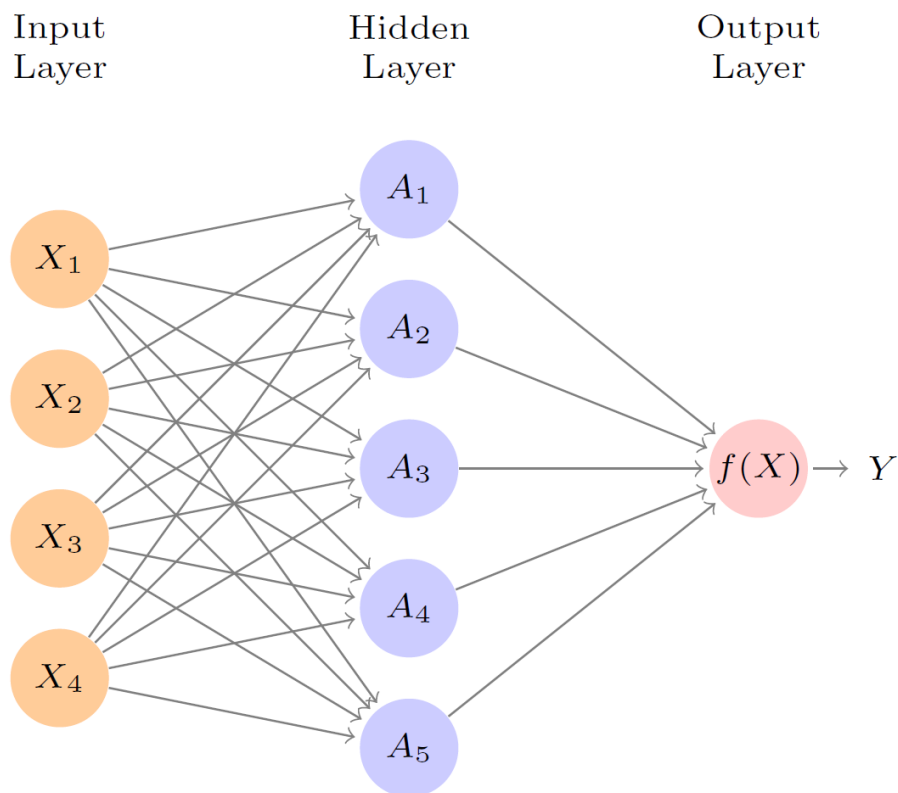$$f(x_i) = \beta_0 + \sum_{k=1}^{K} \beta_k g\left(w_{k0} + \sum_{j=1}^{p} w_{kj} x_{ij}\right)$$

➢ $p = \#inputs, K = \#activations\,(nodes\ in\ the\ hidden\ layer)$

Number of parameters $= (p + 1) \times K + (K + 1) = pK + 2K + 1$

Example: $p = 4, K = 5 \Rightarrow pK + 2K + 1 = 31$

Issues in Fitting and Solution

➢ Nonconvex problem: local vs. global minimum

➢ Solution: slow learning using gradient descent

Issues in Fitting and Solution

➢ Large number of parameters: overfitting

➢ Regularization

  ➢ Ridge or Lasso regularization

$$R(\theta; \lambda) = -\sum_{i=1}^{n} \sum_{m=0}^{9} y_{im} \log(f_m(x_i)) + \lambda \sum_j \theta_j^2$$

  ➢ Dropout learning: randomly remove a fraction of units

# More Complex Neural Networks

➢ Handwritten digit recognition problem

Every image has $p = 28 \times 28 = 784$ pixels.
Each pixel has a number (0-255) to represent its darkness.
Output $Y = (Y_0, Y_1, ..., Y_9)$ where $Y_i = 1, Y_{j, j \neq i} = 0$ if class label $= i$.

➤ L_1 has 256 units, L_2 has 128 units

Input layer

$$\begin{aligned} A_k^{(1)} &= h_k^{(1)}(X) \\ &= g\left(w_{k0}^{(1)} + \sum_{j=1}^{p} w_{kj}^{(1)} X_j\right) \end{aligned}$$

Hidden layer L_1

Hidden layer L_2

Output layer

$$\begin{aligned} Z_m &= \beta_{m0} + \sum_{\ell=1}^{K_2} \beta_{m\ell} h_\ell^{(2)}(X) \\ &= \beta_{m0} + \sum_{\ell=1}^{K_2} \beta_{m\ell} A_\ell^{(2)}, \\ f_m(X) = \Pr(Y = m|X) &= \frac{e^{Z_m}}{\sum_{\ell=0}^{9} e^{Z_\ell}} \end{aligned}$$

X_1

X_2

X_3

X_4

X_5

X_6

$\vdots$

X_p

$A_1^{(1)}$

$A_2^{(1)}$

$A_3^{(1)}$

$A_4^{(1)}$

$\vdots$

$A_{K_1}^{(1)}$

$A_1^{(2)}$

$A_2^{(2)}$

$A_3^{(2)}$

$\vdots$

$A_{K_2}^{(2)}$

$f_0(X) \rightarrow Y_0$

$f_1(X) \rightarrow Y_1$

$\vdots \rightarrow \vdots$

$f_9(X) \rightarrow Y_9$

$\mathbf{W}_1$

$\mathbf{W}_2$

$\mathbf{B}$

$$\begin{aligned} A_\ell^{(2)} &= h_\ell^{(2)}(X) \\ &= g\left(w_{\ell 0}^{(2)} + \sum_{k=1}^{K_1} w_{\ell k}^{(2)} A_k^{(1)}\right) \end{aligned}$$

Number of Parameters

➢ L_1: $785 \times 256 = 200{,}960$

➢ L_2: $257 \times 128 = 32{,}896$

➢ Output layer: $129 \times 10 = 1290$

➢ Total = 235,146

Performance Comparison

➢ **MNIST** data: 60,000 training images, 10,000 test images

➢ Prediction performance

Method	Test Error
Neural Network + Ridge Regularization	2.3%
Neural Network + Dropout Regularization	1.8%
Multinomial Logistic Regression	7.2%
Linear Discriminant Analysis	12.7%

➤ Image classification problem

Every image has $p = 32 \times 32 = 1024$ pixels.
Each pixel has 3 numbers (0-255) to represent red, green and blue.
Output $Y = (Y_0, Y_1, ...)$ where $Y_i = 1, Y_{j, j \neq i} = 0$ if class label $= i$.

CNN

➢ CNNs mimic how humans classify images, by recognizing specific features or patterns anywhere in the image that distinguish each particular object class.

➢ (1) Identify low-level features in the input image, e.g., small edges, patches of color. (2) Combine them to form higher-level features, e.g., parts of ears and eyes. (3) Presence or absence of these features contributes to the probability of given class.
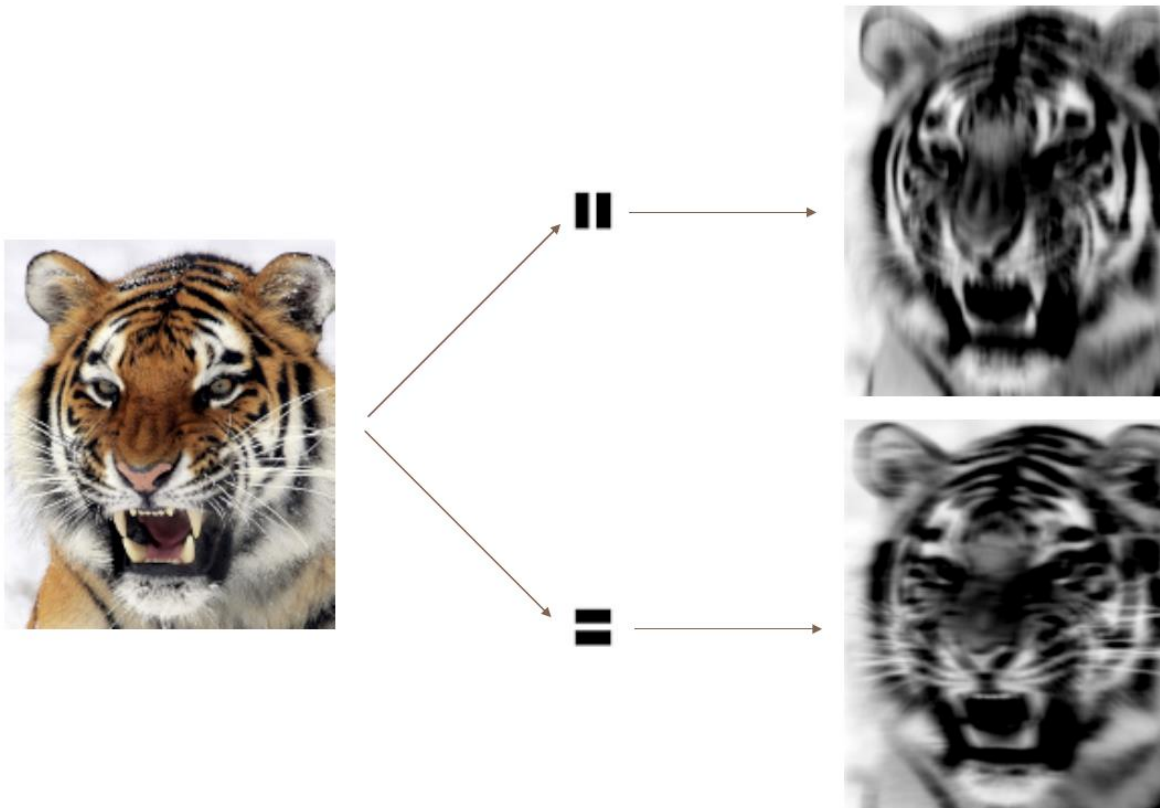
(baamboozle.com)
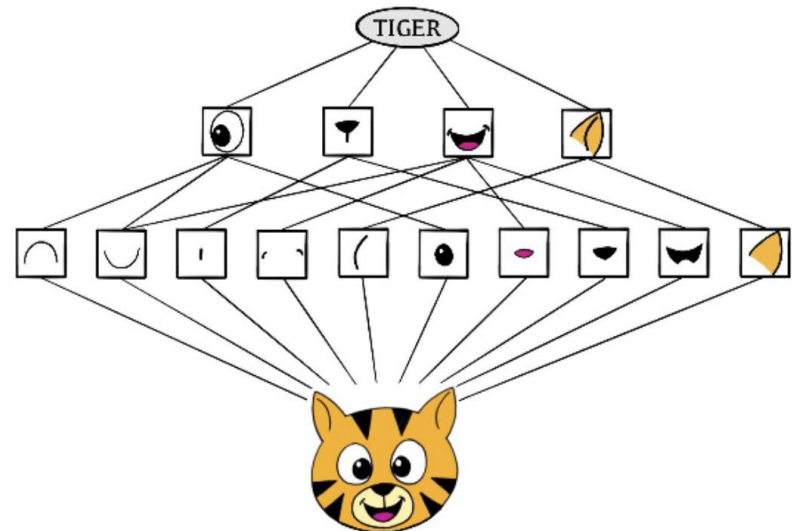
➢ **Convolution layers**: search for small, local patterns in the image using convolution filers

➢ Pooling layers: downsample the small patterns to select a prominent subset

$$\text{Max pool} \begin{bmatrix} 1 & 2 & 5 & 3 \\ 3 & 0 & 1 & 2 \\ 2 & 1 & 3 & 4 \\ 1 & 1 & 2 & 0 \end{bmatrix} \rightarrow \begin{bmatrix} 3 & 5 \\ 2 & 4 \end{bmatrix}$$

➤ A *convolve-then-pool* sequence

➤ **imagenet** dataset: millions of images

flamingo		Cooper's hawk		Cooper's hawk	
flamingo	0.83	kite	0.60	fountain	0.35
spoonbill	0.17	great grey owl	0.09	nail	0.12
white stork	0.00	robin	0.06	hook	0.07
Lhasa Apso		cat		Cape weaver	
Tibetan terrier	0.56	Old English sheepdog	0.82	jacamar	0.28
Lhasa	0.32	Shih-Tzu	0.04	macaw	0.12
cocker spaniel	0.03	Persian cat	0.04	robin	0.12

3. Recurrent Neural Network (RNN)

➢ Many data sources are sequential in nature, calling for special treatment when building predictive models.

  ➢ Documents such as book and movie reviews, articles, tweets.

  ➢ Time series of temperature, rainfall, wind speed, air quality.

  ➢ Financial times series such as mark indices, trading volumes, stock and bond prices, and exchange rates.

  ➢ Recorded speech, musical recordings, and other sound recordings.

  ➢ Handwriting, such as doctor's notes and handwritten zip codes.

➢ RNNs are designed to accommodate and take advantage of the sequential nature of such data.

Example

➤ Document classification problem

Online Movie Review

This has to be one of the worst films of the 1990s. When my friends & I were watching this film (being the target audience it was aimed at) we just sat & watched the first half an hour with our jaws touching the floor at how bad it really was. The rest of the time, everyone else in the theater just started talking to each other, leaving or generally crying into their popcorn …

Output/response: sentiment of the review (positive or negative)

➢ **Input:** $X = \{X_1, X_2, \dots, X_L\}$ is a document with a sequence of L words.

➢ **Output:** often a scalar like the binary sentiment label

$$A_{\ell k} = g\left(w_{k0} + \sum_{j=1}^{p} w_{kj} X_{\ell j} + \sum_{s=1}^{K} u_{ks} A_{\ell-1,s}\right) \qquad O_{\ell} = \beta_0 + \sum_{k=1}^{K} \beta_k A_{\ell k}$$

Application 1: Sentiment Analysis of Movie Reviews

➢ **IMDb** dataset: short critiques of movies

➢ A RNN with 32 hidden units

➢ Trained with dropout regularization on 25,000 reviews

➢ Prediction accuracy = 76%

➢ Using more elaborate versions (LSTM: *long term and short term* memory), performance improved to 87%

➤ **NYSE** dataset: historical trading statistics from the New York Stock Exchange during Dec 1962~Dec 1986

Log trading volume: fraction of all outstanding shares traded on a day

Dow Jones return: difference between log (Dow Jones Industrial Index) on consecutive trading days

**Log volatility**: absolute values of daily price movements

Prediction Problem

➢ Predicting stock price is a notoriously hard problem, but predicting **trading volume** based on recent past history is more manageable (useful for planning trading strategies).

➢ We wish to predict a value of trading volume v_t using

  ➢ Past values of trading volume: $v_{t-1}, v_{t-2}, ...$
  ➢ Past values of Dow Jones return: $r_{t-1}, r_{t-2}, ...$
  ➢ Past values of volatility: $z_{t-1}, z_{t-2}, ...$

➢ Input $X = \{X_1, X_2, ..., X_L\}$ and output Y

$$X_1 = \begin{pmatrix} v_{t-L} \\ r_{t-L} \\ z_{t-L} \end{pmatrix}, \quad X_2 = \begin{pmatrix} v_{t-L+1} \\ r_{t-L+1} \\ z_{t-L+1} \end{pmatrix}, \cdots, X_L = \begin{pmatrix} v_{t-1} \\ r_{t-1} \\ z_{t-1} \end{pmatrix}, \quad \text{and } Y = v_t$$

L is a predefined length (called *lag* in the context of time series)

Performance

➢ Training data (Dec 1962 ~ Jan 1980)

➢ $L = 5$

➢ 12 hidden units in the fitted RNN

➢ $R^2 = 0.42$ (The forecasted series accounts for 42% of the variance of the true data)

When to Use Deep Learning

Success Stories of Deep Learning

➢ CNNs in image classification: machine diagnosis of mammograms, digital X-ray images, ophthalmology eye scans, annotations of MRI scans

➢ RNNs in speech and language translation, forecasting, and document modeling.

➢ **Question:** Should we discard all the older tools, and use deep learning on every problem with data?

Model Comparison

➢ **Salary** dataset: predicting salary of a baseball player in 1987 using his performance statistics in 1986. $p = 19, n = 263$.

➢ Split into a training set (176 players) and test set (87 players)

➢ Comparison of three methods

Model	# Parameters	Mean Abs. Error	Test Set R^2
Linear Regression	20	254.7	0.56
Lasso	12	252.3	0.51
Neural Network	1409	257.4	0.54

➢ Performance vs. complexity tradeoff

➢ Deep learning is attractive when sample size of training data is extremely high, and interpretability is not a high priority.