

SDSC 3006 L02

Class 9. SVM and PCA

Name: Yiren Liu

Email: yirenliu2-c@my.cityu.edu.hk

School of Data Science
City University of Hong Kong

Outline

- **Support Vector Machine**
- **Principal Components Analysis**
- **Principal Components Regression**

Support Vector Machine

Implementation

Notice:

1. SVM is an extension of the support vector classifier, using different kernels(non-linear).
2. Use the e1071 library to demonstrate the SVM on a two-dimensional example.

```
library(e1071)
##Generate data with nonlinear boundary
set.seed(1)
x=matrix(rnorm(200*2),ncol=2)
x[1:100,]=x[1:100,]+2
x[101:150,]=x[101:150,]-2
y=c(rep(1,150),rep(2,50))
dat=data.frame(x=x,y=as.factor(y))
plot(x,col=(3-y))
```

Implementation

##Use SVM with a radial kernel

```
train=sample(200,100)
```

```
svmfit=svm(y~.,data=dat[train,],kernel="radial",gamma=1,cost=1)
```

##"cost" is similar to tuning parameter C, but with opposite

##effects: small "cost", wide margin; large "cost", narrow margin

```
plot(svmfit,dat[train,])
```

```
summary(svmfit)
```

Implementation

##Select best values for "gamma" and "cost" by CV

```
tune.out=tune(svm,y~.,data=dat[train,],kernel="radial",ranges  
=list(cost=c(0.1,1,10,100,1000),gamma=c(0.5,1,2,3,4)))  
summary(tune.out)
```

##Use SVM with a polynomial kernel

```
svmfit=svm(y~.,data=dat[train,],kernel="polynomial" ,  
degree=2,ranges=list(cost=c(0.001,0.01,0.1,1,5,10,100)))  
plot(svmfit,dat[train,])
```

Principal Components Analysis

Details

- Objective:

$$\underset{\phi_{11}, \dots, \phi_{p1}}{\text{maximize}} \left\{ \frac{1}{n} \sum_{i=1}^n \left(\sum_{j=1}^p \phi_{j1} x_{ij} \right)^2 \right\}$$

$$\text{subject to } \sum_{j=1}^p \phi_{j1}^2 = 1$$

- Notice: The principal components are the eigenvectors of a covariance matrix, and hence they are orthogonal. Importantly, the dataset on which PCA technique is to be used must be **scaled**. The results are also sensitive to the relative scaling.

Details


- The solution of PCA also could treat as:

$$Au = \lambda u$$

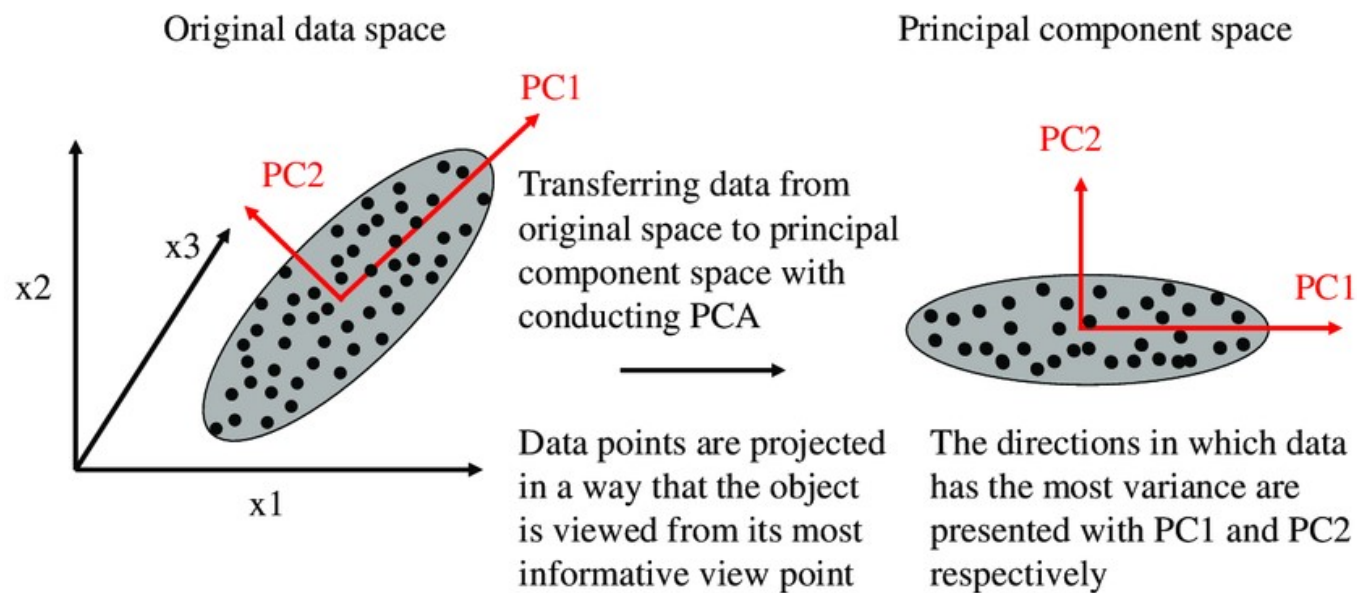
- Where $A \in R^{p \times p}$, it is the covariance matrix of dataset X . u and λ are eigen vector and eigen value respectively.
- It can easily show that eigen vectors are **orthogonal** with each other.

Examples

Example 1:
2D case:



Example 2:
3D case:



PC1 will find the direction in which the data has the most variance. In this case, the points vary most in the direction of PC1.

Implementation

Notice:

1. We perform PCA on the USArrests data set, which is part of the base R package.
2. Using the `prcomp()` function, which is one of several functions in R that perform PCA.
3. By default, the `prcomp()` function centers the variables to have mean zero. 'Option `scale = TRUE`' scale the variables to have sd 1.

Implementation

```
names (USArrests)
```

```
pr.out = prcomp (USArrests , scale = TRUE)
```

```
names (pr.out)
```

```
##rotation: principal component loadings
```

```
pr.out$rotation
```

```
biplot (pr.out , scale = 0)
```

```
##sdev: standard deviation of principal components
```

```
pr.out$sdev
```

```
pr.var = pr.out$sdev^2
```

Implementation

```
##compute the proportion of variance explained by each principal component
```

```
pve = pr.var / sum (pr.var)
```

```
par (mfrow = c(1, 2))
```

```
plot (pve , xlab = " Principal Component ", ylab = " Proportion of Variance Explained ",  
ylim = c(0, 1), type = "b")
```

```
plot ( cumsum (pve), xlab = " Principal Component ", ylab = " Cumulative Proportion of  
Variance Explained ", ylim = c(0, 1), type = "b")
```

Principal Components Regression

Details

- The principal components regression (PCR) approach involves constructing the first M principal components, Z_1, \dots, Z_M , and then using these components as the predictors in a linear regression model that is fit using least squares.

Details

- If the assumption underlying PCR holds, then fitting a least squares model to Z_1, \dots, Z_M will lead to better results than fitting a least squares model to X_1, \dots, X_p , since most or all of the information in the data that relates to the response is contained in Z_1, \dots, Z_M , and by estimating only $M \ll p$ coefficients we can mitigate overfitting.

Implementation

Notice:

1. PCR can be performed using `pcr()` function in `pls` library
2. We now apply PCR to the `Hitters` data, in order to predict Salary.

```
install.packages('pls')  
library(ISLR)  
library(pls)  
attach(Hitters)  
Hitters = na.omit (Hitters) ##remove rows with missing values  
set.seed (2)  
pcr.fit = pcr(Salary ~ ., data = Hitters , scale = TRUE, validation = "CV")  
##10-fold cross-validation  
summary (pcr.fit)  
##plot the CV scores  
validationplot(pcr.fit , val.type = "MSEP")
```

Implementation

```
##create training set and test set
```

```
x=model.matrix(Salary~.,Hitters)[-1]
```

```
y=Hitters$Salary
```

```
train=sample(1:nrow(x), nrow(x)/2)
```

```
test=(-train)
```

```
y.test=y[test]
```

```
##perform PCR on training set
```

```
set.seed(1)
```

```
pcr.fit = pcr(Salary ~ ., data = Hitters , subset = train , scale = TRUE , validation = "CV")
```

```
validationplot (pcr.fit , val.type = "MSEP") ##find the best M
```

```
##compute test MSE
```

```
pcr.pred = predict(pcr.fit , x[test , ], ncomp = 5)
```

```
mean ((pcr.pred - y.test)^2) ##compare with shrinkage method?
```

```
##perform PCR on full data
```

```
pcr.fit = pcr (y ~ x, scale = TRUE , ncomp = 5)
```